

# davos: a Python package “smuggler” for constructing lightweight reproducible notebooks

Paxton C. Fitzpatrick, Jeremy R. Manning\*

*Department of Psychological and Brain Sciences  
Dartmouth College, Hanover, NH 03755*

---

## Abstract

Reproducibility is a core requirement of modern scientific research. For computational research, reproducibility means that code should produce the same results, even when run on different systems. A standard approach to ensuring reproducibility entails packaging a project’s dependencies along with its primary code base. Existing solutions vary in how deeply these dependencies are specified, ranging from virtual environments, to containers, to virtual machines. Each of these existing solutions requires installing or setting up a system for running the desired code, increasing the complexity and time cost of sharing or engaging with reproducible science. Here, we propose a lighter-weight solution: the **davos** library. When used in combination with a notebook-based Python project, **davos** provides a mechanism for specifying (and automatically installing) the correct versions of the project’s dependencies. The **davos** library further ensures that those packages and specific versions are used every time the notebook’s code is executed. This enables researchers to share a complete reproducible copy of their code within a single Jupyter notebook file.

*Keywords:* Reproducibility, Open science, Python, Jupyter Notebook, Google Colaboratory, Package management

---

---

\*Corresponding author

*Email address:* `Jeremy.R.Manning@Dartmouth.edu` (Jeremy R. Manning)

## Required Metadata

### Current code version

Nr.	Code metadata description	Metadata value
C1	Current code version	v0.1.1
C2	Permanent link to code/repository used for this code version	<a href="https://github.com/ContextLab/davos/tree/v0.1.1">https://github.com/ContextLab/davos/tree/v0.1.1</a>
C3	Code Ocean compute capsule	
C4	Legal Code License	MIT
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python, JavaScript, PyPI/pip, IPython, Jupyter, Ipykernel, PyZMQ. Additional tools used for tests: pytest, Selenium, Requests, mypy, GitHub Actions
C7	Compilation requirements, operating environments, and dependencies	Dependencies: Python $\geq 3.6$ , packaging, setuptools. Supported OSes: MacOS, Linux, Unix-like. Supported IPython environments: Jupyter notebooks, JupyterLab, Google Colaboratory, Binder, IDE-based notebook editors.
C8	Link to developer documentation/manual	<a href="https://github.com/ContextLab/davos#readme">https://github.com/ContextLab/davos#readme</a>
C9	Support email for questions	contextualdynamics@gmail.com

Table 1: Code metadata

## 1. Motivation and significance

The same computer code may not behave identically under different circumstances. For example, when code depends on external libraries, different versions of those libraries may function differently. Or when CPU or GPU instruction sets differ across machines, the same high-level code may be compiled into different machine instructions. Because executing identical code does not guarantee identical outcomes, code sharing alone is often insufficient for enabling researchers to reproduce each other’s work, or to collaborate on projects involving data collection or analysis.

Within the Python [1] community, external packages that are published in the most popular repositories [2, 3] are associated with version numbers and

12 tags that allow users to guarantee they are installing exactly the same code  
13 across different computing environments [4]. While it is *possible* to manually  
14 install the intended version of every dependency of a Python script or pack-  
15 age, manually tracking down those dependencies can impose a substantial  
16 burden on the user and create room for mistakes and inconsistencies. Fur-  
17 ther, when dependency versions are left unspecified, replicating the original  
18 computing environment becomes difficult or impossible.

19 Computational researchers and other programmers have developed a broad  
20 set of approaches and tools to facilitate code sharing and reproducible out-  
21 comes (Fig. 1). At one extreme, simply distributing a set of Python scripts  
22 (.py files) may enable others to use or gain insights into the relevant work.  
23 Because Python is installed by default on most modern operating systems,  
24 for some projects, this may be sufficient. Another popular approach entails  
25 creating Jupyter notebooks [8] that comprise a mix of text, executable code,  
26 and embedded media. Notebooks may call or import external scripts or  
27 libraries—even intersperse snippets of other programming or markup lang-  
28 uages—in order to provide a more compact and readable experience for users.  
29 Both of these systems (Python scripts and notebooks) provide a convenient  
30 means of sharing code, with the caveat that they do not specify the comput-  
31 ing environment in which the code is executed. Therefore the functionality of  
32 code shared using these systems cannot be guaranteed across different users  
33 or setups.

34 At another extreme, virtual machines [9, 10, 11] provide a hardware-level  
35 simulation of the desired system. Virtual machines are typically isolated such  
36 that installing or running software on a virtual machine does not impact the  
37 user’s primary operating system or computing environment. Containers [e.g.,  
38 12, 13] provide a similar “isolated” experience. Although containerized envi-  
39 ronments do not specify hardware-level operations, they are typically pack-  
40 aged with a complete operating system, in addition to a complete copy of  
41 Python and any relevant package dependencies. Virtual environments [e.g.,  
42 6, 7] also provide a computing environment that is largely separated from the  
43 user’s main environment. They incorporate a copy of Python and the target  
44 software’s dependencies, but virtual environments do not specify or repro-  
45 duce an operating system for the runtime environment. Each of these systems  
46 (virtual machines, containers, and virtual environments) guarantees (to dif-  
47 fering degrees—at the hardware level, operating system level, and Python  
48 environment level, respectively) that the relevant code will run similarly for  
49 different users. However, each of these systems also relies on additional soft-  
50 ware that can be complex or resource-intensive to install and use, creating  
51 potential barriers to both contributing to and taking advantage of open sci-  
52 ence resources.

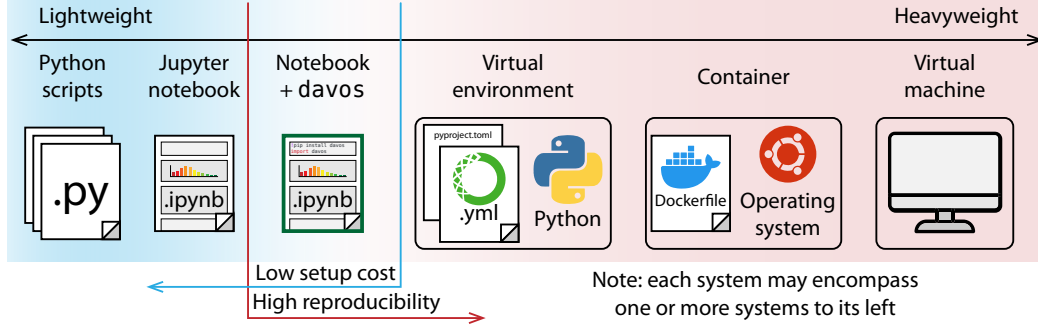


Figure 1: **Systems for sharing code within the Python ecosystem.** From left to right: plain-text **Python scripts** (.py files) provide the most basic “system” for sharing raw code. Scripts may reference external libraries, but those libraries must be manually installed on other users’ systems. Further, any checking needed to verify that the correct versions of those libraries were installed must also be performed manually. **Jupyter notebooks** (.ipynb files) comprise embedded text, executable code, and media (including rendered figures, code output, etc.). When the **davos** library is imported into a Jupyter notebook, the notebook’s functionality is extended to automatically install any required external libraries (at their correct versions, when specified). **Virtual environments** allow users to install an isolated copy of Python and all required dependencies. This typically entails distributing a configuration file (e.g., a `pyproject.toml` [5] or `environment.yml`) that specifies all project dependencies (including version numbers of external libraries) alongside the primary code base. Users can then install a third-party tool [e.g., 6, 7] to read the file and build the environment. **Containers** provide a means of defining an isolated environment that includes a complete operating system (independent of the user’s operating system), in addition to (optionally) specifying a virtual environment or other configurations needed to provide the necessary computing environment. Containers are typically defined using specification files (e.g., a plain-text `Dockerfile`) that instruct the virtualization engine regarding how to build the containerized environment. **Virtual machines** provide a complete hardware-level simulation of the computing environment. In addition to simulating specific hardware, virtual machines (typically specified using binary images files) must also define operating system-level properties of the computing environment. Systems to the left of the blue vertical line entail sharing individual files, with no additional installation or configuration needed to run the target code. Systems to the right of the red vertical line support precise control over dependencies and versioning. Notebooks enhanced using the **davos** library are easily shareable and require minimal setup costs, while also facilitating high reproducibility by enabling precise control over project dependencies.

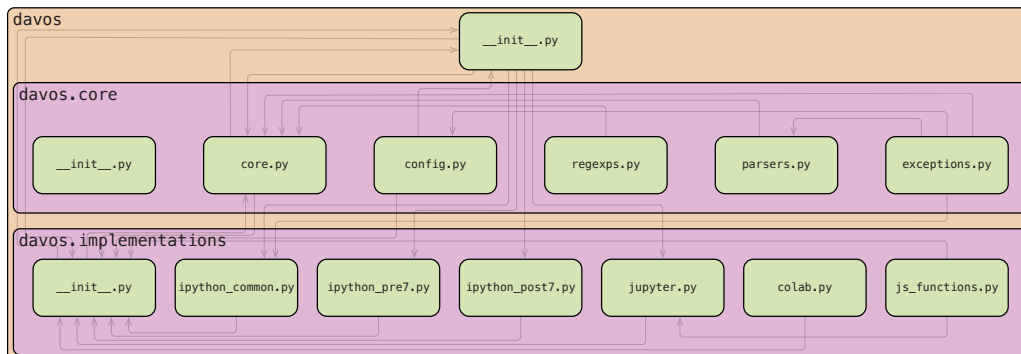


Figure 2: **Package structure.**

We designed **davos** to occupy a “sweet spot” between these extremes. **davos** is a notebook-installable package that adds functionality to the default notebook experience. Like standard Jupyter notebooks, **davos**-enhanced notebooks allow researchers to include text, executable code, and media within a single file. No further setup or installation is required, beyond what is needed to run standard Jupyter notebooks. And like virtual environments, **davos** provides a convenient mechanism for fully specifying (and installing, as needed) a complete set of Python dependencies, including package versions.

## 2. Software description

The **davos** package is named after Davos Seaworth, a smuggler often referred to as “the Onion Knight” from the series *A Song of Ice and Fire* by George R. R. Martin.

### 2.1. Software architecture

The **davos** package consists of two interdependent subpackages (see Fig. 2). The first, **davos.core**, comprises a set of modules that implement the bulk of the package’s core functionality, including pipelines for installing and validating packages, custom parsers for the **smuggle** statement (see Section 2.2.1) and onion comment (see Section 2.2.2), and a runtime interface for configuring **davos**’s behavior (see Section 2.2.3). However, certain critical aspects of this functionality require (often substantially) different implementations depending on properties of the notebook environment in which **davos** is used (e.g., whether the frontend is provided by Jupyter or Google Colaboratory, or which version of IPython [14] is used by the notebook kernel). To deal with this, environment-dependent parts of core features and behaviors are isolated and abstracted to “helper functions” in the **davos.implementations**

subpackage. This second subpackage defines multiple, interchangeable versions of each helper function, organized into modules by the conditions that trigger their use. At runtime, `davos` detects various features in the notebook environment and selectively imports a single version of each helper function into the top-level `davos.implementations` namespace, allowing `davos.core` modules to access the proper implementations for the current notebook environment in a single, consistent location. An additional benefit of this design pattern is that it allows both maintainers and users to easily extend `davos` to support new, updated, or custom notebook variants by creating a new `davos.implementations` module that defines its own version of each helper function, modified as much or little as necessary.

## 2.2. Software functionalities

### 2.2.1. The *smuggle* statement

Functionally, importing `davos` in an IPython notebook enables an additional Python keyword: “`smuggle`” (see Section 2.3 for details on how this works). The `smuggle` statement can be used as a drop-in replacement for Python’s built-in “`import`” to load libraries, modules, and other objects into the current namespace. However, whereas `import` will fail if the requested package is not installed locally, `smuggle` statements can handle missing packages on the fly. If a smuggled package does not exist in the local environment, `davos` will install it automatically, expose its contents to Python’s `import` machinery, and load it into the namespace for immediate use.

### 2.2.2. The *onion* comment

For greater control over the behavior of `smuggle` statements, `davos` defines an additional construct called the “onion comment.” An onion comment is a special type of inline comment that may be placed on a line containing a `smuggle` statement to customize how `davos` searches for the smuggled package locally and, if necessary, downloads and installs it. Onion comments follow a simple format based on the “type comment” syntax introduced in PEP 484 [15], and are designed to make managing packages with `davos` intuitive and familiar. To construct an onion comment, users provide the name of the installer program (e.g., `pip`) and the same arguments one would use to manually install the package as desired via the command line:

```

# enable smuggle statements
import davos

# if numpy is not installed locally, pip-install it and display verbose output
smuggle numpy as np          # pip: numpy --verbose

# pip-install pandas without using or writing to the package cache
smuggle pandas as pd         # pip: pandas --no-cache-dir

# install scipy from a relative local path, in editable mode
from scipy.stats smuggle ttest_ind  # pip: -e ../../pkgs/scipy

```

Occasionally, a package's distribution name (i.e., the name used when installing it) may differ from its top-level module name (i.e., the name used when importing it). In such cases, an onion comment can be used to ensure that **davos** installs the proper distribution if the smuggled package can't be found locally:

```

# package is named "python-dateutil" on PyPI, but imported as "dateutil"
smuggle dateutil          # pip: python-dateutil

# package is named "scikit-learn" on PyPI, but imported as "sklearn"
from sklearn.decomposition smuggle PCA  # pip: scikit-learn

```

However, the most powerful use of the onion comment is making **smuggle** statements *version-sensitive*. If an onion comment includes a version specifier [4], **davos** will ensure that the version of the package loaded into the notebook matches the specific version requested, or satisfies the given version constraints. If the smuggled package exists locally, **davos** will extract its version info from its metadata and compare it to the specifier provided. If the two are incompatible (or no local installation is found), **davos** will install and load a suitable version of the package instead:

```

# specifically use matplotlib v3.4.2, pip-installing it if needed
smuggle matplotlib.pyplot as plt  # pip: matplotlib==3.4.2

# use a version of seaborn no older than v0.9.1, but prior to v0.11
smuggle seaborn as sns          # pip: seaborn>=0.9.1,<0.11

```

Onion comments can also be used to smuggle specific VCS references (e.g., Git [16] branches, commits, tags, etc.):

```

# use quail as the package existed on GitHub at commit 6c847a4
smuggle quail  # pip: git+https://github.com/ContextLab/quail.git@6c847a4

```

130 **davos** processes onion comments internally before forwarding arguments to  
131 the installer program. In addition to preventing onion comments from being  
132 used as a vehicle for shell injection attacks, this enables **davos** to take certain  
133 logical actions when particular arguments are passed. For example, each of  
134 the `-I/--ignore-installed`, `-U/--upgrade`, and `--force-reinstall` flags  
135 will cause **davos** to skip searching for a smuggled package locally before  
136 installing a new copy:

```
# install hypertools v0.7 without first checking for it locally
smuggle hypertools as hyp          # pip: hypertools==0.7 --ignore-installed

# always install the latest version of requests, including pre-releases
from requests smuggle Session     # pip: requests --upgrade --pre
```

138 Similarly, passing `--no-input` will temporarily enable **davos**'s non-interactive  
139 mode (see Section 2.2.3), and installing a smuggled package into `<dir>` with  
140 `--target <dir>` will cause **davos** to prepend `dir` to the module search path  
141 (i.e., `sys.path`), if necessary, so the package can be imported.

### 142 2.2.3. The *davos* config

143 The **davos** config provides a simple, high-level interface for controlling  
144 various aspects of **davos**'s behavior. After importing **davos**, the `davos.config`  
145 object (a singleton) exposes a number of configurable options as attributes  
146 that can be assigned different values, checked at runtime, and displayed in  
147 the notebook cell output (see Fig. 4 for example usage). These include:

- 148 • `.active`: This option allows users to disable **davos** functionality (i.e.,  
149 support for **smuggle** statements and onion comments) for subsequent  
150 notebook cells by setting its value to `False`. **davos** can be re-enabled  
151 at any time by setting this option to `True` (the default when **davos** is  
152 first imported). See Section 2.3 for additional info.
- 153 • `.auto_rerun`: This option controls how **davos** behaves when attempt-  
154 ing to **smuggle** a new version of a package that was previously imported  
155 and cannot be reloaded. This can happen if the package includes extension  
156 modules that dynamically link C or C++ objects to the Python  
157 interpreter itself, and the code that generates those objects was changed  
158 between the old and new versions. If this option is set to `True`, **davos**  
159 will automatically restart the notebook kernel and rerun all code up to  
160 (and including) the current **smuggle** statement. If `False` (the default),  
161 **davos** will instead issue a warning, pause execution, and prompt the  
162 user with buttons to either restart and rerun the notebook, or con-  
163 tinue running with the previously imported package version until the



164 next time the kernel is restarted manually. (Note: not configurable in  
165 Google Colaboratory).

- 166 • `.confirm_install`: If `True` (default: `False`), `davos` will require user  
167 confirmation (`[y]es/[n]o` input) before installing a smuggled package.
- 168 • `.noninteractive`: Setting this option to `True` (default: `False`) en-  
169 ables non-interactive mode, in which all user input and confirmation  
170 is disabled. Note that in non-interactive mode, the `confirm_install`  
171 option is set to `False`, and if `auto_rerun` is `False`, `davos` will throw an  
172 error if a smuggled package cannot be reloaded, rather than prompting  
173 the user.
- 174 • `.pip_executable`: This option's value specifies the path to the `pip`  
175 executable used to install smuggled packages. The default is program-  
176 matically determined from the Python environment and falls back to  
177 `sys.executable -m pip` if no executable can be found.
- 178 • `.suppress_stdout`: If `True` (default: `False`), suppress all unnecessary  
179 output issued by both `davos` and the installer program. This can be  
180 useful when smuggling packages that need to install many dependencies  
181 and/or generate extensive output. If the installer program throws an  
182 error, both the stdout and stderr streams will be displayed with the  
183 traceback.

184 The top-level `davos` namespace also defines a handful of convenience func-  
185 tions for setting and checking `davos`'s active/inactive state (`davos.activate()`;  
186 `davos.deactivate()`; `davos.is_active()`) as well as the `davos.configure()`  
187 function, which allows setting multiple config options at once.

### 188 2.3. Implementation details

189 Functionally, importing `davos` appears to make “`smuggle`” a valid Python  
190 keyword, similar to standard keywords like “`import`”, “`def`”, or “`return`”.  
191 It also appears to fundamentally change how Python treats comments, sud-  
192 denly allowing them to potentially influence code behavior at runtime if they  
193 conform to a particular syntax. However, `davos` doesn't actually modify the  
194 rules of Python's parser or lexical analyzer in order to accomplish this—  
195 in fact, modifying the Python grammar isn't possible at runtime, as doing  
196 so would require rebuilding the interpreter. Instead, `davos` leverages the  
197 IPython notebook backend to implement the `smuggle` statement and onion  
198 comment via a combination of namespace injections and its own (far simpler)  
199 custom parser.

200 The `smuggle` keyword can be enabled and disabled at any time by “ac-  
201 tivating” and “deactivating” `davos` (see Section 2.2.3, above). When `davos`  
202 is first imported, it is activated automatically. Activating `davos` triggers  
203 two actions: (1) the `smuggle()` function is injected into the IPython user  
204 namespace, and (2) the `davos` parser is registered as a custom IPython input  
205 transformer. IPython preprocesses all executed code as plain text before it is  
206 sent to the Python compiler, in order to handle special constructs like `%magic`  
207 and `!shell` commands. `davos` hooks into this process to transform `smuggle`  
208 statements into syntactically valid Python code. The `davos` parser uses a  
209 complex regular expression [17] to match lines of code containing `smuggle`  
210 statements (and, optionally, onion comments), extract relevant information  
211 from their text, and replace them with equivalent calls to the `smuggle()`  
212 function. For example, if a user runs a notebook cell containing

```
213 smuggle numpy as np      # pip: numpy>1.16,<=1.20 -vv
```

214 the code that is actually executed by the Python interpreter would be

```
215 smuggle(name="numpy", as_="np", installer="pip",  
        args_str="\"numpy>1.16,<=1.20 -vv\"",  
        installer_kwargs={'editable': False,  
                          'spec': 'numpy>1.16,<=1.20',  
                          'verbosity': 2})
```

216 The call to the `smuggle()` function then carries out `davos`’s central logic  
217 of determining whether the smuggled package should be installed, doing so  
218 if necessary, and subsequently loading it into the namespace. This process  
219 is outlined in Figure 3. Because the `smuggle()` function is defined in the  
220 notebook namespace, it is also possible (though never necessary) to call  
221 it directly. Deactivating `davos` will delete the name “`smuggle`” from the  
222 namespace, unless its value has been overwritten and no longer refers to the  
223 `smuggle()` function. It will also deregister the `davos` parser from the set of  
224 input transformers run when each notebook cell is executed. While the over-  
225 head added by the `davos` parser is minimal, this may be useful, for example,  
226 when optimizing or precisely profiling code.

### 227 3. Illustrative Examples

228 Figure 4 illustrates how one might use `davos` in a real-world setting to  
229 facilitate reproducible data analyses. In this scenario, a user has acquired a  
230 dataset and pre-trained semantic model (possibly shared by a collaborator,  
231 or downloaded from a public repository) with which they want to run a set

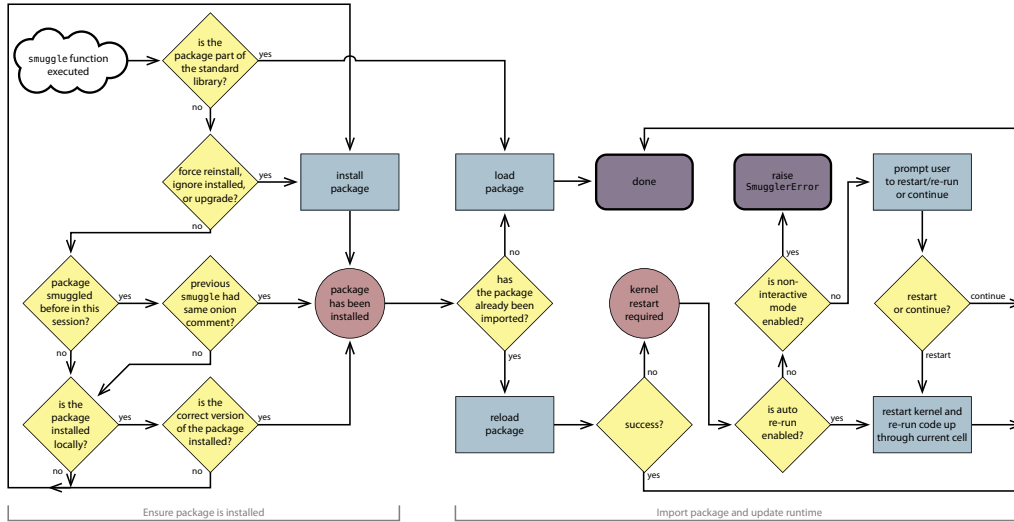


Figure 3: `smuggle()` function logic.

of analyses in a Jupyter notebook. The hypothetical dataset (a `pandas` [18] Panel) and model (a `scikit-learn` [19] Pipeline) were created and saved using the latest versions of each package available in June 2019 and, due to their age and the formats in which they were provided, now pose certain challenges to loading and using them. The code in Figure 4 would be included at the top of the user’s analysis notebook, and demonstrates how `davos` enables them to overcome these issues cleanly and efficiently, while simultaneously constructing a reproducible Python environment that ensures their analyses are always run under the same conditions, both by the user themselves and anyone with whom they share their code.

After importing `davos` (line 1), the user first smuggles two utilities for interacting with local files in the code below. The `smuggle` statement in line 3 loads the `is_file()` function from the Python standard library’s `os.path` module. Standard library modules are included with all Python distributions, so this line is functionally equivalent to an `import` statement and does not need or benefit from an onion comment. Line 4 loads the `joblib` library [20], installing it first, if necessary. Since `joblib`’s I/O interface has historically remained stable and backwards-compatible across releases, requiring that users have a particular exact version installed would likely be unnecessarily restrictive. However, there’s also no guarantee that a *future* release won’t introduce some breaking change, so to help ensure the analysis notebook continues to run properly in the future, the onion comment in line 4 limits allowable versions to those already released when the code was written.

Line 6 then uses the `davos.config` object to enable `davos`’s `auto_re-`

```

1  import davos
2
3  from os.path smuggle is_file
4  smuggle joblib                                # pip: joblib<=1.2.0
5
6  davos.config.auto_rerun = True
7  smuggle numpy as np                            # pip: numpy==1.21.6
8
9  if not is_file("~/datasets/data-new.csv"):
10     smuggle pandas as pd                        # pip: pandas<0.25.0
11     tmp_data = pd.read_pickle("~/datasets/data-old.pkl")
12     tmp_data.to_frame().to_csv("~/datasets/data-new.csv")
13
14  smuggle pandas as pd                          # pip: pandas==1.3.5
15
16  davos.configure(auto_rerun=False, suppress_stdout=True, noninteractive=True)
17  smuggle tensorflow as tf                      # pip: tensorflow==2.9.2
18  from umap smuggle UMAP                      # pip: umap-learn[plot,parametric_umap]==0.5.3
19  davos.configure(suppress_stdout=False, noninteractive=False)
20
21  smuggle matplotlib.pyplot as plt             # pip: matplotlib==3.5.3
22  smuggle seaborn as sns                      # pip: seaborn==0.12.1
23  smuggle quail                                # pip: git+https://github.com/myfork/quail@6c847a4
24
25  davos.config.pip_executable = "~/envs/nb-server/bin/pip"
26  smuggle widgetsnbextension as _             # pip: widgetsnbextension==3.5.2
27  davos.config.pip_executable = "~/envs/nb-kernel/bin/pip"
28  smuggle ipywidgets                          # pip: ipywidgets==7.6.5
29
30  from tqdm.notebook smuggle tqdm              # pip: tqdm==4.62.3
31
32  data = pd.read_csv("~/datasets/data-new.csv", index_col=[0, 1])
33  smuggle sklearn                             # pip: scikit-learn<0.22.0
34  transformer = joblib.load("~/models/text-transformer.joblib")
35  smuggle sklearn                             # pip: scikit-learn==1.1.3

```

Figure 4: Example use case for davos's functionality.

run option before smuggling the next two packages: NumPy [21] and pandas. Because these libraries rely heavily on custom C data types, loading the particular versions from the onion comments may require restarting the notebook kernel if different versions were already imported during the same interpreter session (see Section 2.2.3). This is unlikely to happen with NumPy, as it is loaded explicitly only once (by line 7) and is not required by any packages loaded before it (and therefore not imported and cached while loading them). However, it's possible a future user could edit this notebook to add additional code before line 7, run the notebook's cells out of order, or even configure their IPython kernel to execute custom code on startup (e.g., as Google Colab notebooks do), any of which could potentially import a different, existing version of NumPy before davos smuggles the version specified in the onion comment. Enabling the `auto_rerun` option before smuggling NumPy will help to simplify running the notebook in these scenarios, but the main benefit it affords comes when smuggling the pandas library.

The dataset that the user wants to analyze was provided as a pandas.Panel object, serialized with Python's built-in pickle module. The pickle protocol is a popular choice for persisting data in Python, allowing users to save, share, and load arbitrary objects, though with an important caveat: In order to successfully "unpickle" (i.e., load and restore) a "pickled" object, its class must be defined in and importable from the same module as when it was saved. The Panel class was removed from pandas in July, 2019 (v0.25.0), meaning that the dataset can be loaded only if the user's notebook uses a pandas version released prior to that change. However, more recent updates to the pandas library since then have brought substantial improvements that the user would like to take advantage of in their analysis code, including faster performance, more memory-efficient data types, new functions and methods, compatibility with newer versions of Python and popular plotting libraries, and better display formats for data structures (including in Jupyter notebooks). davos makes it possible to do both within a single notebook, and with the `auto_rerun` option, this process can be fully automated. The first time the notebook is run, lines 10–12 will install and load an older version of pandas that defines the Panel class, use it to read the dataset from its serialized file format, convert the dataset from a Panel to a multi-index DataFrame, and write its contents to a CSV file. The `smuggle` statement on line 14 will then install a newer pandas version (v1.3.5) for use in the analyses themselves. Because an older pandas version was previously loaded (by line 10), and because the package's C extension modules were modified between the old and new versions, switching to the new version will require restarting the notebook kernel. When davos determines this, it will automatically restart the kernel, re-run the notebook up through the

current cell, and also re-queue any subsequent cells that were queued before the restart was triggered. This second time the code is run (as well as in any future runs), lines 10–12 will be skipped since the CSV-formatted dataset has already been created, and the `smuggle` statements up through line 14 will execute (virtually) as fast as regular `import` statements since the required package versions have already been installed.

Line 16 then uses the `davos.configure()` function to disable the `auto_rerun` option and simultaneously enable two other options: `suppress_stdout` and `noninteractive`. With these options enabled, lines 17–18 `smuggle TensorFlow` [22], a powerful end-to-end platform for building and working with machine learning models, and `UMAP` [23], a library that implements a family of related manifold learning techniques. The onion comment in line 18 also specifies that `UMAP` should be installed with the optional requirements needed for its “plot” and “parametric\_umap” features. Together, these two packages depend on 36 other unique packages, most of which have dependencies of their own. And if many of these are not already installed in the user’s environment, lines 17–18 could take multiple minutes to run. Enabling the `noninteractive` option ensures the installation process won’t require any input from the user to finish successfully, allowing them to simply ignore the notebook during this time and either work in another browser tab or application, or step away from their computer entirely. Additionally, installing more than three dozen packages could result in the installer program writing an extremely large amount of standard output text to the notebook cell’s output area. Enabling `suppress_stdout` hides the output from these two `smuggle` statements so that other potentially important output (e.g., from previous `smuggle` statements) isn’t buried and lost. Importantly, if either line fails for any reason, both the stdout and stderr streams will be shown alongside the regular Python traceback.

After re-enabling these two options (line 19), the user next smuggles specific versions of three plotting libraries: `Matplotlib` [24], `seaborn` [25], and `Quail` [26] (lines 21–23). Because the first two are requirements of `UMAP`’s optional “plot” feature, they will have already been installed by line 18, though possibly as different versions than those specified in the onion comments on lines 21 and 22. If the installed and specified versions are the same, these `smuggle` statements will function like regular `import` statements to load the libraries into the notebook namespace. If they differ, `davos` will download the requested versions in place of the installed versions before doing so. It’s also worth noting that while loading `UMAP` in line 18 will not result in either `Matplotlib` or `seaborn` being imported internally due to the specific layout of the `UMAP` package, if it had, `davos` would also gracefully replace the cached versions of the two packages so that the smuggled versions are

338 loaded properly. And (as with all other `smuggle` statements in Figure 4), if  
339 for any reason the specified package versions are incompatible with any other  
340 packages (e.g., `UMAP`), the installer program would display a warning to this  
341 effect in the notebook.

342 The third plotting library smuggled, `Quail`, is also pinned to a precise  
343 version with an onion comment, though in a slightly different manner than  
344 other packages in Figure 4. Line 23 installs and loads `Quail` from the user’s  
345 fork (“myfork”) of the package’s GitHub repository, at a specific point in its  
346 revision history (i.e., the commit whose short hash is “6c847a4”). This can  
347 serve as a useful way to access a version of a package the user has customized,  
348 patched, or somehow modified for this specific use case, while still allowing  
349 them to further modify their fork of the package in the future without affect-  
350 ing the version used in these analyses.

351 Next, the user smuggles a pair of packages that extend the Jupyter note-  
352 book interface with various interactive JavaScript widgets. The `ipywid-`  
353 `gets` [27] package provides an API for creating these widgets with Python  
354 code, and the `widgetsnbextension` package provides the machinery for the  
355 notebook frontend to display them. Thus, `ipywidgets` must be installed in  
356 the same environment as the IPython kernel, and `widgetsnbextension` must  
357 be installed in the environment that houses the Jupyter notebook server. In  
358 the simplest case, these two environments are the same; however, a common  
359 approach to using Jupyter notebooks with multiple different Python envi-  
360 ronments (e.g., for different tasks or projects) is to run the notebook server  
361 from a “base environment,” with additional environments each providing  
362 their own separate, interchangeable IPython kernels. (While possible to de-  
363 tect programmatically, Figure 4 assumes this setup to simplify the example  
364 code.) To handle this multi-environment installation, line 25 temporarily sets  
365 the `pip` executable used to install smuggled packages to that of the notebook  
366 server’s environment before smuggling `widgetsnbextension` (line 26). Since  
367 there is no need to actually load this package for later use in the analysis  
368 code, it is aliased as a single underscore—by convention, a “throwaway” vari-  
369 able in Python. In line 27, the user changes the `pip` executable back to that  
370 of the notebook kernel environment (its original value) before installing `ipy-`  
371 `widgets` and subsequent packages. With these two packages installed, line 30  
372 smuggles `tqdm` [28], which provides convenient progress bars for long-running  
373 code. In Jupyter notebooks, the `tqdm.notebook` module can be imported to  
374 enable prettier progress bars displayed via `ipywidgets`, if that package is  
375 installed and importable. Therefore, to take advantage of this feature, it  
376 was important to `smuggle tqdm` after ensuring the `ipywidgets` package was  
377 available.

378 Finally, the user loads in the re-formatted dataset (line 32) and pre-

379 trained model (line 34) they'll use in their analyses. But in accessing the  
 380 latter, the user encounters an issue similar to the one they initially faced  
 381 with the dataset. This hypothetical semantic model was provided as a  
 382 `scikit-learn` "Pipeline" object that allows the user to pass text data  
 383 through two pre-trained transformers, in series: a `CountVectorizer` instance  
 384 and a `LatentDirichletAllocation` instance (i.e., a topic model [29]). The  
 385 model was saved by its original creator using the `joblib` library, as `scikit-`  
 386 `learn`'s documentation recommends, with the caveat that because `joblib`  
 387 uses the `pickle` protocol internally, the ability to save and load pre-trained  
 388 models is not guaranteed across versions due to `pickle`'s requirement that  
 389 referenced classes be importable from the same location when saved and  
 390 loaded. This issue affects the user's model by way of the `LatentDirich-`  
 391 `letAllocation` class—when the Pipeline object was created (again, in  
 392 June, 2019, using `scikit-learn` v0.21.3), the class's definition lived in the  
 393 `sklearn.decomposition.online_lda` module. However, in version 0.22.0  
 394 (released November, 2019), that module was renamed to "`_online_lda`,"  
 395 and in v0.22.1 (January, 2020) it was again renamed to "`_lda`" (which has  
 396 remained its name since then). Thus, in order to successfully load the model  
 397 that includes the pre-trained `LatentDirichletAllocation` instance, in line  
 398 33, the user first smuggles a version of `scikit-learn` prior to v0.22.0 (i.e.,  
 399 before the first time the module's name was changed). Once the model  
 400 is loaded and reconstructed in memory from a compatible package version  
 401 (line 34), the user then upgrades to a newer version of `scikit-learn` for  
 402 use elsewhere in their analyses. When line 35 is run, `davos` will replace the  
 403 existing `scikit-learn` version with v1.1.3 in both the user's local package  
 404 environment and in memory without interruption, such that "`sklearn`" will  
 405 reference the new version in any code below. However, this will not affect the  
 406 already compiled code object stored in the "`transformer`" variable, meaning  
 407 that the pre-trained model will be usable simultaneously.

## 408 4. Impact

409 Like virtual environments, containers, and virtual machines, the `davos`  
 410 library (when used in conjunction with Jupyter notebooks) provides a light-  
 411 weight mechanism for sharing code and ensuring reproducibility across users  
 412 and computing environments (Fig. 1). Further, `davos` enables users to fully  
 413 specify (and install, as needed) any project dependencies within the same  
 414 notebook. This provides a system whereby executable code (along with text  
 415 and media) *and* code for setting up and configuring the project dependencies,  
 416 may be combined within a single notebook file.

417 We designed `davos` for use in research applications. For example, in many



418 settings, `davos` may be used as a drop-in replacement for more-difficult-to-  
419 set-up virtual environments, containers, and/or virtual machines. For re-  
420 searchers, this lowers barriers to both sharing code. By eliminating most  
421 of the setup costs of reconstructing the original researchers’ computing en-  
422 vironment, `davos` also lowers barriers to entry for members of the scientific  
423 community and the public who seek to benefit from shared code.

424 Beyond research applications, `davos` is also useful in pedagogical settings.  
425 For example, in programming courses, instructors and students may import  
426 the `davos` library into their notebooks to provide a simple means of ensur-  
427 ing their code will run on others’ machines. When combined with online  
428 notebook-based platforms like Google Colaboratory, `davos` provides a conve-  
429 nient way to manage dependencies within a notebook, without requiring any  
430 software (beyond a web browser) to be installed on the students’ or instruc-  
431 tors’ systems. For the same reasons, `davos` also provides an elegant means of  
432 sharing ready-to-run notebook-based demonstrations or tutorials that install  
433 their dependencies automatically.

434 Since its initial release, `davos` has found use in a variety of applications.  
435 In addition to managing computing environments for multiple ongoing re-  
436 search studies, `davos` is being used by both students and instructors in pro-  
437 gramming and methods courses such as Storytelling with Data [30] (an open  
438 course on data science, visualization, and communication) and Laboratory  
439 in Psychological Science [31] (an open course on experimental and statistical  
440 methods for psychology research) to simplify distributing lessons and sub-  
441 mitting assignments, as well as in online demos such as `abstract2paper` [32]  
442 (an example application of GPT-Neo [33, 34]) to share ready-to-run code  
443 that installs dependencies automatically.

444 Our work also has several more subtle “advanced” use cases and potential  
445 impacts. Whereas Python’s built-in `import` statement is agnostic to pack-  
446 ages’ version information, `smuggle` statements (when combined with onion  
447 comments) are version-sensitive. And because onion comments are parsed  
448 at runtime, required packages and their specified versions are installed in a  
449 just-in-time manner. Thus, it is possible in most cases to `smuggle` a specific  
450 package version or revision even if a different version has already been loaded.  
451 This enables more complex uses that take advantage of multiple versions of  
452 a package within a single interpreter session. This could be useful in cases  
453 where specific features are added or removed from a package across differ-  
454 ent versions, or in comparing the performance or functionality of particular  
455 features across different versions of the same package.

456 A second advanced use case is in providing a proof-of-concept of how one  
457 can add new “keyword-like” operators to the Python language by leverag-  
458 ing notebooks’ error-handling mechanisms. This could lead to exciting new

459 tools that, like `davos`, extend the Python language in useful ways within  
460 notebook-based environments. We note that our approach to adding the  
461 `smuggle` keyword to Python when `davos` is imported into a notebook-based  
462 environment also has the potential to be exploited for more nefarious pur-  
463 poses. For example, a malicious user could use a similar approach (e.g.,  
464 in a different library) to substantially change a notebook’s functionality by  
465 adding new *unexpected* keyword-like objects (e.g., based around common ty-  
466 pos). This could lead to difficult-to-predict changes in a notebook’s behavior  
467 once the malicious library was imported. This highlights an important rea-  
468 son why security-conscious users would be well-served to only make use of  
469 libraries from trusted sources, or whose code is publicly available for review.

## 470 5. Conclusions

471 The `davos` library supports reproducible research by providing a novel  
472 lightweight system for sharing notebook-based code. But perhaps the most  
473 exciting uses of the `davos` library are those that we have *not* yet considered  
474 or imagined. We hope that the Python community will find `davos` to pro-  
475 vide a convenient means of managing project dependencies to facilitate code  
476 sharing. We also hope that some of the more advanced applications of our  
477 library might lead to new insights or discoveries.

## 478 Author Contributions

479 **Paxton C. Fitzpatrick:** Conceptualization, Methodology, Software,  
480 Validation, Writing - Original Draft, Visualization. **Jeremy R. Manning:**  
481 Conceptualization, Resources, Validation, Writing - Review & Editing, Su-  
482 pervision, Funding acquisition.

## 483 Funding

484 Our work was supported in part by NSF grant number 2145172 to JRM.  
485 The content is solely the responsibility of the authors and does not necessarily  
486 represent the official views of our supporting organizations.

## 487 Declaration of Competing Interest

488 We wish to confirm that there are no known conflicts of interest associated  
489 with this publication and there has been no significant financial support for  
490 this work that could have influenced its outcome.

## 491 Acknowledgements

492 We acknowledge useful feedback and discussion from the students of  
493 JRM's *Storytelling with Data* course (Winter, 2022 offering) who used pre-  
494 liminary versions of our library in several assignments.

## 495 References

- 496 [1] G. van Rossum, Python reference manual, Department of Computer  
497 Science [CS] (R 9525) (1995).
- 498 [2] Python Software Foundation, The Python Package Index (PyPI),  
499 <https://pypi.org> (2003).
- 500 [3] conda-forge community, The conda-forge Project: Community-based  
501 Software Distribution Built on the conda Package Format and Ecosys-  
502 tem, <https://doi.org/10.5281/zenodo.4774217> (July 2015). doi:  
503 [10.5281/zenodo.4774217](https://doi.org/10.5281/zenodo.4774217).
- 504 [4] N. Coghlan, D. Stufft, Version Identification and Dependency Specifica-  
505 tion, PEP 440, Python Software Foundation (March 2013).
- 506 [5] B. Cannon, N. Smith, D. Stufft, Specifying minimum build system re-  
507 quirements for python projects, PEP 518, Python Software Foundation  
508 (May 2016).
- 509 [6] Anaconda, Inc., conda, <https://docs.conda.io> (2012).
- 510 [7] S. Eustace, Poetry: Python packaging and dependency management  
511 made easy, <https://github.com/python-poetry/poetry> (December  
512 2019).
- 513 [8] T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier,  
514 J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov,  
515 D. Avila, S. Abdalla, C. Willing, Jupyter Notebooks – a publish-  
516 ing format for reproducible computational workflows, in: F. Loizides,  
517 B. Schmidt (Eds.), Positioning and Power in Academic Publishing: Play-  
518 ers, Agents and Agendas, IOS Press, Netherlands, 2016, pp. 87–90.  
519 doi:[10.3233/978-1-61499-649-1-87](https://doi.org/10.3233/978-1-61499-649-1-87).
- 520 [9] R. P. Goldberg, Survey of virtual machine research, Computer 7 (6)  
521 (1974) 34–45.

- [10] Y. Altintas, C. Brecher, M. Weck, S. Witt, Virtual Machine Tool, CIRP Annals 54 (2) (2005) 115–138. doi:[https://doi.org/10.1016/S0007-8506\(07\)60022-5](https://doi.org/10.1016/S0007-8506(07)60022-5).
- [11] M. Rosenblum, VMware’s Virtual Platform: A virtual machine monitor for commodity PCs, in: IEEE Hot Chips Symposium, IEEE, 1999, pp. 185–196.
- [12] D. Merkel, Docker: lightweight linux containers for consistent development and deployment, Linux Journal 239 (2) (2014) 2.
- [13] G. M. Kurtzer, V. Sochat, M. W. Bauer, Singularity: Scientific containers for mobility of compute, PLoS One 12 (5) (2017) e0177459.
- [14] F. Pérez, B. E. Granger, IPython: a system for interactive scientific computing, Computing in science and engineering 9 (3) (2007) 21–29. doi:[10.1109/MCSE.2007.53](https://doi.org/10.1109/MCSE.2007.53).
- [15] G. van Rossum, J. Lehtosalo, L. Langa, Type Hints, PEP 484, Python Software Foundation (September 2014).
- [16] L. Torvalds, J. Hamano, Git: Fast version control system, <https://git.kernel.org/pub/scm/git/git.git> (April 2005).
- [17] K. Thompson, Programming Techniques: Regular expression search algorithm, Communications of the ACM 11 (6) (1968) 419–422. doi:[10.1145/363347.363387](https://doi.org/10.1145/363347.363387).
- [18] W. McKinney, Data Structures for Statistical Computing in Python, in: S. van der Walt, J. Millman (Eds.), Proceedings of the 9th Python in Science Conference, 2010, pp. 56–61. doi:[10.25080/Majors-92bf1922-00a](https://doi.org/10.25080/Majors-92bf1922-00a).
- [19] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: machine learning in Python, Journal of Machine Learning Research 12 (2011) 2825–2830.
- [20] G. Varoquaux, Joblib: Computing with Python functions, <https://github.com/joblib/joblib> (July 2010).
- [21] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith,

- 555 R. Kern, M. Pícus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane,  
556 J. F. del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Shep-  
557 pard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, T. E. Oliphant,  
558 Array programming with NumPy, *Nature* 585 (7825) (2020) 357–362.  
559 [doi:10.1038/s41586-020-2649-2](https://doi.org/10.1038/s41586-020-2649-2).
- 560 [22] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S.  
561 Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow,  
562 A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kud-  
563 lur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah,  
564 M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker,  
565 V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wat-  
566 tenberg, M. Wicke, Y. Yu, X. Zheng, [TensorFlow: Large-Scale Ma-](https://www.tensorflow.org/)  
567 [chine Learning on Heterogeneous Systems](https://www.tensorflow.org/), software available from ten-  
568 sorflow.org (2015).  
569 URL <https://www.tensorflow.org/>
- 570 [23] L. McInnes, J. Healy, N. Saul, L. Großberger, UMAP: Uniform Manifold  
571 Approximation and Projection, *Journal of Open Source Software* 3 (29)  
572 (2018) 861. [doi:https://doi.org/10.21105/joss.00861](https://doi.org/10.21105/joss.00861).
- 573 [24] J. D. Hunter, Matplotlib: A 2D graphics environment, *Computing in*  
574 *Science and Engineering* 9 (3) (2007) 90–95. [doi:10.1109/MCSE.2007.](https://doi.org/10.1109/MCSE.2007.55)  
575 [55](https://doi.org/10.1109/MCSE.2007.55).
- 576 [25] M. L. Waskom, seaborn: statistical data visualization, *Journal of Open*  
577 *Source Software* 6 (60) (2021) 3021. [doi:10.21105/joss.03021](https://doi.org/10.21105/joss.03021).
- 578 [26] A. C. Heusser, P. C. Fitzpatrick, C. E. Field, K. Ziman, J. R. Manning,  
579 Quail: a Python toolbox for analyzing and plotting free recall data,  
580 *Journal of Open Source Software* 10.21105/joss.00424 (2017).
- 581 [27] J. Frederic, J. Grout, Jupyter Widgets Contributors, ipywidgets: In-  
582 teractive Widgets for the Jupyter Notebook, [https://github.com/](https://github.com/jupyter-widgets/ipywidgets)  
583 [jupyter-widgets/ipywidgets](https://github.com/jupyter-widgets/ipywidgets) (August 2015).
- 584 [28] C. da Costa-Luis, S. K. Larroque, K. Altendorf, H. Mary, richardsheri-  
585 dan, M. Korobov, N. Raphael, I. Ivanov, M. Bargull, N. Rodrigues,  
586 G. Chen, A. Lee, C. Newey, CrazyPython, JC, M. Zugnoni, M. D.  
587 Pagel, mjestevens777, M. Dektyarev, A. Rothberg, A. Plavin, D. Pan-  
588 teleit, F. Dill, FichteFoll, G. Sturm, HeoHeo, H. van Kemenade, J. Mc-  
589 Cracken, MapleCCC, M. Nordlund, tqdm: A Fast, Extensible Progress

- 590 Bar for Python and CLI, <https://github.com/tqdm/tqdm> (September  
591 2022). doi:10.5281/zenodo.595120.
- 592 [29] D. M. Blei, A. Y. Ng, M. I. Jordan, Latent dirichlet allocation, Journal  
593 of Machine Learning Research 3 (2003) 993–1022.
- 594 [30] J. R. Manning, Storytelling with Data, [https://github.com/](https://github.com/ContextLab/storytelling-with-data)  
595 [ContextLab/storytelling-with-data](https://github.com/ContextLab/storytelling-with-data) (June 2021). doi:10.5281/  
596 [zenodo.5182775](https://github.com/ContextLab/storytelling-with-data).
- 597 [31] J. Manning, ContextLab/experimental-psychology: v1.0 (Spring, 2022),  
598 [https://github.com/ContextLab/experimental-psychology/tree/](https://github.com/ContextLab/experimental-psychology/tree/v1.0)  
599 [v1.0](https://github.com/ContextLab/experimental-psychology/tree/v1.0) (May 2022). doi:10.5281/zenodo.6596762.
- 600 [32] J. R. Manning, abstract2paper, [https://github.com/ContextLab/](https://github.com/ContextLab/abstract2paper)  
601 [abstract2paper](https://github.com/ContextLab/abstract2paper) (June 2021). doi:10.5281/zenodo.7261831.
- 602 [33] L. Gao, S. Biderman, S. Black, L. Golding, T. Hoppe, C. Foster,  
603 J. Phang, H. He, A. Thite, N. Nabeshima, S. Presser, C. Leahy, The  
604 Pile: An 800GB Dataset of Diverse Text for Language Modeling, arXiv  
605 preprint arXiv:2101.00027 (2020).
- 606 [34] S. Black, L. Gao, P. Wang, C. Leahy, S. Biderman, GPT-Neo: Large  
607 Scale Autoregressive Language Modeling with Mesh-Tensorflow, [http:](http://github.com/eleutherai/gpt-neo)  
608 [//github.com/eleutherai/gpt-neo](http://github.com/eleutherai/gpt-neo) (2021).