

davos: a Python package “smuggler” for constructing lightweight reproducible notebooks

Paxton C. Fitzpatrick, Jeremy R. Manning*

*Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755*

Abstract

Reproducibility is a core requirement of modern scientific research. For computational research, reproducibility means that code should produce the same results, even when run on different systems. A standard approach to ensuring reproducibility entails packaging a project’s dependencies along with its primary code base. Existing solutions vary in how deeply these dependencies are specified, ranging from virtual environments, to containers, to virtual machines. Each of these existing solutions requires installing or setting up a system for running the desired code, increasing the complexity and time cost of sharing or engaging with reproducible science. Here, we propose a lighter-weight solution: the **davos** library. When used in combination with a notebook-based Python project, **davos** provides a mechanism for specifying (and automatically installing) the correct versions of the project’s dependencies. The **davos** library further ensures that those packages and specific versions are used every time the notebook’s code is executed. This enables researchers to share a complete reproducible copy of their code within a single Jupyter notebook file.

Keywords: Reproducibility, Open science, Python, Jupyter Notebook, Google Colaboratory, Package management

*Corresponding author

Email address: `Jeremy.R.Manning@Dartmouth.edu` (Jeremy R. Manning)

Required Metadata

Current code version

Nr.	Code metadata description	Metadata value
C1	Current code version	v0.1.1
C2	Permanent link to code/repository used for this code version	https://github.com/ContextLab/davos/tree/v0.1.1
C3	Code Ocean compute capsule	
C4	Legal Code License	MIT
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python, JavaScript, PyPI/pip, IPython, Jupyter, Ipykernel, PyZMQ. Additional tools used for tests: pytest, Selenium, Requests, mypy, GitHub Actions
C7	Compilation requirements, operating environments, and dependencies	Dependencies: Python ≥ 3.6 , packaging, setuptools. Supported OSes: MacOS, Linux, Unix-like. Supported IPython environments: Jupyter notebooks, JupyterLab, Google Colaboratory, Binder, IDE-based notebook editors.
C8	Link to developer documentation/manual	https://github.com/ContextLab/davos#readme
C9	Support email for questions	contextualdynamics@gmail.com

Table 1: Code metadata

1. Motivation and significance

The same computer code may not behave identically under different circumstances. For example, when code depends on external libraries, different versions of those libraries may function differently. Or when CPU or GPU instruction sets differ across machines, the same high-level code may be compiled into different machine instructions. Because executing identical code does not guarantee identical outcomes, code sharing alone is often insufficient for enabling researchers to reproduce each other’s work, or to collaborate on projects involving data collection or analysis.

Within the Python [1] community, external packages that are published in the most popular repositories [2, 3] are associated with version numbers and

12 tags that allow users to guarantee they are installing exactly the same code
13 across different computing environments [4]. While it is *possible* to manually
14 install the intended version of every dependency of a Python script or pack-
15 age, manually tracking down those dependencies can impose a substantial
16 burden on the user and create room for mistakes and inconsistencies. Fur-
17 ther, when dependency versions are left unspecified, replicating the original
18 computing environment becomes difficult or impossible.

19 Computational researchers and other programmers have developed a broad
20 set of approaches and tools to facilitate code sharing and reproducible out-
21 comes (Fig. 1). At one extreme, simply distributing a set of Python scripts
22 (.py files) may enable others to use or gain insights into the relevant work.
23 Because Python is installed by default on most modern operating systems,
24 for some projects, this may be sufficient. Another popular approach entails
25 creating Jupyter notebooks [8] that comprise a mix of text, executable code,
26 and embedded media. Notebooks may call or import external scripts or
27 libraries—even intersperse snippets of other programming or markup lang-
28 uages—in order to provide a more compact and readable experience for users.
29 Both of these systems (Python scripts and notebooks) provide a convenient
30 means of sharing code, with the caveat that they do not specify the comput-
31 ing environment in which the code is executed. Therefore the functionality of
32 code shared using these systems cannot be guaranteed across different users
33 or setups.

34 At another extreme, virtual machines [9, 10, 11] provide a hardware-level
35 simulation of the desired system. Virtual machines are typically isolated such
36 that installing or running software on a virtual machine does not impact the
37 user’s primary operating system or computing environment. Containers [e.g.,
38 12, 13] provide a similar “isolated” experience. Although containerized envi-
39 ronments do not specify hardware-level operations, they are typically pack-
40 aged with a complete operating system, in addition to a complete copy of
41 Python and any relevant package dependencies. Virtual environments [e.g.,
42 6, 7] also provide a computing environment that is largely separated from the
43 user’s main environment. They incorporate a copy of Python and the target
44 software’s dependencies, but virtual environments do not specify or repro-
45 duce an operating system for the runtime environment. Each of these systems
46 (virtual machines, containers, and virtual environments) guarantees (to dif-
47 fering degrees—at the hardware level, operating system level, and Python
48 environment level, respectively) that the relevant code will run similarly for
49 different users. However, each of these systems also relies on additional soft-
50 ware that can be complex or resource-intensive to install and use, creating
51 potential barriers to both contributing to and taking advantage of open sci-
52 ence resources.

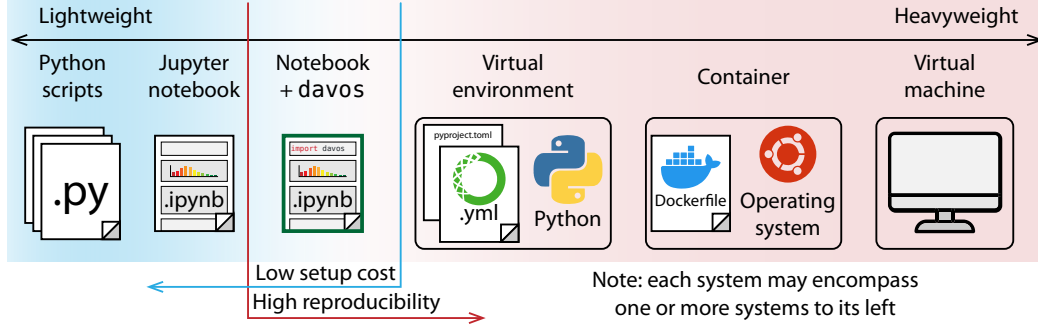


Figure 1: **Systems for sharing code within the Python ecosystem.** From left to right: plain-text **Python scripts** (.py files) provide the most basic “system” for sharing raw code. Scripts may reference external libraries, but those libraries must be manually installed on other users’ systems. Further, any checking needed to verify that the correct versions of those libraries were installed must also be performed manually. **Jupyter notebooks** (.ipynb files) comprise embedded text, executable code, and media (including rendered figures, code output, etc.). When the **davos** library is imported into a Jupyter notebook, the notebook’s functionality is extended to automatically install any required external libraries (at their correct versions, when specified). **Virtual environments** allow users to install an isolated copy of Python and all required dependencies. This typically entails distributing a configuration file (e.g., a `pyproject.toml` [5] or `environment.yml`) that specifies all project dependencies (including version numbers of external libraries) alongside the primary code base. Users can then install a third-party tool [e.g., 6, 7] to read the file and build the environment. **Containers** provide a means of defining an isolated environment that includes a complete operating system (independent of the user’s operating system), in addition to (optionally) specifying a virtual environment or other configurations needed to provide the necessary computing environment. Containers are typically defined using specification files (e.g., a plain-text `Dockerfile`) that instruct the virtualization engine regarding how to build the containerized environment. **Virtual machines** provide a complete hardware-level simulation of the computing environment. In addition to simulating specific hardware, virtual machines (typically specified using binary images files) must also define operating system-level properties of the computing environment. Systems to the left of the blue vertical line entail sharing individual files, with no additional installation or configuration needed to run the target code. Systems to the right of the red vertical line support precise control over dependencies and versioning. Notebooks enhanced using the **davos** library are easily shareable and require minimal setup costs, while also facilitating high reproducibility by enabling precise control over project dependencies.

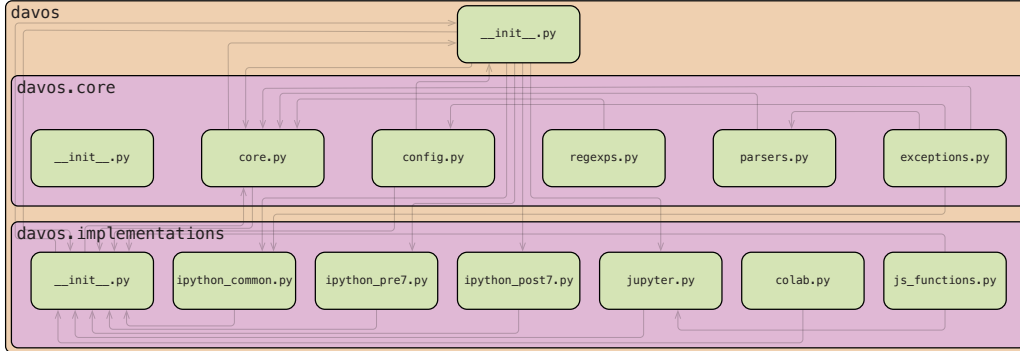


Figure 2: **Package structure.**

We designed **davos** to occupy a “sweet spot” between these extremes. **davos** is a notebook-installable package that adds functionality to the default notebook experience. Like standard Jupyter notebooks, **davos**-enhanced notebooks allow researchers to include text, executable code, and media within a single file. No further setup or installation is required, beyond what is needed to run standard Jupyter notebooks. And like virtual environments, **davos** provides a convenient mechanism for fully specifying (and installing, as needed) a complete set of Python dependencies, including package versions.

2. Software description

2.1. Software architecture

The **davos** package consists of two interdependent subpackages (see Fig. 2). The first, **davos.core**, comprises a set of modules that implement the bulk of the package’s core functionality, including pipelines for installing and validating packages, custom parsers for the **smuggle** statement (see Section 2.2.1) and onion comment (see Section 2.2.2), and a runtime interface for configuring **davos**’s behavior (see Section 2.2.3). However, certain critical aspects of this functionality require (often substantially) different implementation approaches depending on various properties of the notebook environment in which **davos** is used (e.g., whether the frontend is provided by Jupyter or Google Colaboratory, or which version of IPython [14] is used by the notebook kernel). To deal with this, environment-dependent parts of core features and behaviors are isolated and abstracted to “helper functions” in the **davos.implementations** subpackage. This second subpackage defines multiple, interchangeable versions of each helper function, organized into modules by the conditions that trigger their use. At runtime, **davos** detects various features in the notebook environment and selectively imports a single

79 version of each helper function into the top-level `davos.implementations`
80 namespace, allowing `davos.core` modules to access the correct implementa-
81 tions for the current notebook environment in a single, consistent location.
82 An additional benefit of this design pattern is that it allows maintainers or
83 users to easily extend `davos` to support new, updated, or custom notebook
84 variants simply by creating a new `davos.implementations` module with any
85 necessary tweaks to existing helper functions.

86 *2.2. Software functionalities*

87 *2.2.1. The `smuggle` statement*

88 Importing `davos` in an IPython notebook enables an additional Python
89 keyword: “`smuggle`” (see Section 2.3 for details on how this works). The
90 `smuggle` statement can be used as a drop-in replacement for Python’s built-
91 in `import` statement to load libraries, modules, and other objects into the
92 current namespace. However, whereas `import` will fail if the requested pack-
93 age is not installed locally, `smuggle` statements can handle missing packages
94 on the fly. If a smuggled package does not exist in the local environment,
95 `davos` will install it automatically, expose its contents to Python’s `import`
96 machinery, and load it into the namespace for immediate use.

97 *2.2.2. The onion comment*

98 For greater control over the behavior of `smuggle` statements, `davos` de-
99 fines an additional construct called the “onion comment”. An onion comment
100 is a special type of inline comment that may be placed on a line containing a
101 `smuggle` statement to customize how `davos` searches for the smuggled pack-
102 age locally and, if necessary, downloads and installs it. Onion comments
103 follow a simple syntax based on the “type comment” syntax introduced in
104 PEP 484 [15], and are designed to make managing packages with `davos` intu-
105 itive and familiar. To construct an onion comment, simply provide the name
106 of the installer program (e.g., `pip`) and the same arguments one would use
107 to manually install the package as desired via the command line:

```
import davos

# if numpy is not installed locally, pip-install it and display verbose output
smuggle numpy as np      # pip: numpy --verbose

# pip-install pandas without using or writing to the package cache
smuggle pandas as pd     # pip: pandas --no-cache-dir

# install scipy from a relative local path, in editable mode
from scipy.stats smuggle ttest_ind      # pip: -e ../../pkgs/scipy
```

108

109 Onion comments are useful when smuggling a package whose distribution
110 name (i.e., the name used when installing it) is different from its top-level
111 module name (i.e., the name used when importing it):

```
112     smuggle dateutil      # pip: python-dateutil
    from sklearn.decomposition smuggle PCA      # pip: scikit-learn
```

113 However, the most powerful use of the onion comment is making `smuggle`
114 statements *version-sensitive*. If an onion comment includes a version spec-
115 ifier [4], `davos` will ensure that the version of the package loaded into the
116 notebook matches the specific version requested, or satisfies the given ver-
117 sion constraints. If the smuggled package exists locally, `davos` will extract
118 its version info from its metadata and compare it to the specifier provided. If
119 the two are incompatible (or no local installation is found), `davos` will install
120 and load a suitable version of the package instead:

```
121     # specifically use matplotlib v3.4.2, pip-installing it if needed
    smuggle matplotlib.pyplot as plt      # pip: matplotlib==3.4.2

    # use a version of seaborn no older than v0.9.1, but before v0.11
    smuggle seaborn as sns      # pip: seaborn>=0.9.1,<0.11
```

122 Onion comments can similarly be used to smuggle specific VCS references
123 (e.g., Git [16] branches, commits, tags, etc.):

```
124     # use quail as the package existed on GitHub at commit 6c847a4
    smuggle quail      # pip: git+https://github.com/ContextLab/quail.git@6c847a4
```

125 `davos` processes onion comments internally before forwarding arguments to
126 the installer program. In addition to preventing onion comments from being
127 used as a vehicle for shell injection attacks, this allows `davos` take certain
128 logical actions when particular arguments are passed. For example, the `-I/-`
129 `-ignore-installed`, `-U/--upgrade`, and `--force-reinstall` flags will all
130 cause `davos` to skip searching for a smuggled package locally before installing
131 a new copy:

```
132     # install hypertools v0.7 without first checking for it locally
    smuggle hypertools as hyp      # pip: hypertools==0.7 --ignore-installed

    # always install the latest version of requests, including pre-releases
    from requests smuggle Session      # pip: requests --upgrade --pre
```

133 Similarly, passing `--no-input` will temporarily enable `davos`'s non-interactive
134 mode (see Section 2.2.2), and installing a smuggled package into `<dir>` with
135 `--target <dir>` will cause `dir` to be prepended to the module search path
136 (`sys.path`), if necessary, so the package can be imported

137 2.2.3. The *davos* config

138 The `davos` config object provides a simple, high-level interface that allows
139 users to view and set various options that affect `davos`'s behavior. After
140 importing `davos`, the config instance (a singleton) for the current session is
141 available as `davos.config`, and its various fields are accessible as attributes.
142 The config object exposes a mixture of writable and read-only fields. Writable
143 fields include:

- 144 • `.active`: Whether or not `davos` functionality (i.e., support for `smug-`
145 `gle` statements and onion comments) should be enabled for subsequent
146 code. Defaults to `True` when `davos` is first imported. See Section 2.3
147 for additional info.
- 148 • `.auto_rerun`: Controls behavior if `davos` is used to `smuggle` a new ver-
149 sion of a package that was previously imported and cannot be reloaded
150 (i.e., it contains C-extensions that dynamically generate code). If `True`
151 (default: `False`), `davos` will automatically restart the notebook kernel
152 and rerun all code up to (and including) the current `smuggle` state-
153 ment. Otherwise, `davos` will issue a warning, pause execution, and
154 prompt the user with buttons to either restart & rerun the notebook
155 or continue running with the imported package version. (Note: not
156 configurable in Google Colaboratory).
- 157 • `.confirm_install`: If `True` (default: `False`), `davos` will require user
158 confirmation (`[y]es/[n]o` input) before installing a smuggled package.
- 159 • `.noninteractive`: Setting to `True` (default: `False`) enables non-interactive
160 mode, in which all user input and confirmation is disabled. Note that
161 in non-interactive mode, the `confirm_install` option is set to `False`,
162 and if `auto_rerun` is `False`, `davos` will throw an error if a smuggled
163 package cannot be reloaded.
- 164 • `.pip_executable`: The path to the `pip` executable used to install
165 smuggled packages. Default is programmatically determined from the
166 Python environment and falls back to `sys.executable -m pip` if one
167 can't be found.

- `.suppress_stdout`: If `True` (default: `False`), suppress all unnecessary output issued by both `davos` and the installer program. Useful when smuggling packages that need to install many dependencies and/or generate extensive output. If the installer program throws an error, both `stdout` and `stderr` will be shown with the traceback.

The top-level `davos` namespace additionally defines a handful of convenience functions for setting and checking `davos`'s active/inactive state (`davos.activate()`; `davos.deactivate()`; `davos.is_active()`) as well as the `davos.configure()` function, which allows setting multiple config fields at once.

2.3. Implementation details

Functionally, importing `davos` appears to define “`smuggle`” as a Python keyword, similar to “`import`”, “`def`”, or “`return`”. It also appears to cause comments to be parsed, and their contents potentially able to affect code behavior, which they normally are not. However, `davos` doesn't actually modify the rules of Python's parser or lexical analyzer—in fact, modifying the Python grammar isn't possible at runtime, as doing so would require rebuilding the interpreter. Instead, `davos` leverages the IPython notebook backend to implement the `smuggle` statement and onion comment via a combination of namespace injections and its own (far simpler) custom parser.

The `smuggle` keyword can be enabled and disabled at any time by “activating” and “deactivating” `davos` (see Section 2.2.3, above). When `davos` is first imported, it is activated automatically. Activating `davos` triggers two actions: (1) the `smuggle()` function is injected into the IPython user namespace, and (2) the `davos` parser is registered as a custom IPython input transformer. IPython preprocesses all executed code as plain text before it is sent to the Python parser, in order to handle special constructs like `%magic` and `!shell` commands. `davos` hooks into this process to transform `smuggle` statements into syntactically valid Python code. The `davos` parser uses a complex regular expression [17] to match lines of code containing `smuggle` statements (and, optionally, onion comments), extract relevant information from their text, and replace them with equivalent calls to the `smuggle()` function. For example, if a user runs a notebook cell containing

```
smuggle numpy as np      # pip: numpy>1.16,<=1.20 -vv
```

the code that is actually executed by the Python interpreter would be

```

smuggle(name="numpy", as_="np", installer="pip",
        args_str="numpy>1.16,<=1.20 -vv",
        installer_kwargs={'editable': False,
                           'spec': 'numpy>1.16,<=1.20',
                           'verbosity': 2})

```

202

203 Because the `smuggle()` function is defined in the notebook namespace, it is
 204 also possible (though never necessary) to call it directly. Deactivating `davos`
 205 will delete the name “`smuggle`” from the namespace, unless its value has
 206 been overwritten and no longer refers to the `smuggle()` function. It will also
 207 deregister the `davos` parser from the set of input transformers run when each
 208 notebook cell is executed. While the overhead added by the `davos` parser is
 209 de minimis, this may be useful, for example, when optimizing or precisely
 210 profiling code.

211 3. Illustrative Examples

212 4. Impact

213 Like virtual environments, containers, and virtual machines, the `davos` li-
 214 brary (when used in conjunction with Jupyter notebooks) provides a lightweight
 215 mechanism for sharing code and ensuring reproducibility across users and
 216 computing environments (Fig. 1). Further, `davos` enables users to fully
 217 specify (and install, as needed) any project dependencies within the same
 218 notebook. This provides a system whereby executable code (along with text
 219 and media) *and* code for setting up and configuring the project dependencies,
 220 may be combined within a single notebook file.

221 We designed `davos` for use in research applications. For example, in many
 222 settings `davos` may be used as a drop-in replacement for more-difficult-to-
 223 set-up virtual environments, containers, and/or virtual machines. For re-
 224 searchers, this lowers barriers to sharing code. By eliminating most of the
 225 setup costs of reconstructing the original researchers’ computing environ-
 226 ment, `davos` also lowers barriers to entry for members of the scientific com-
 227 munity and the public who seek to *benefit* from shared code.

228 Beyond research applications, `davos` is also useful in pedagogical settings.
 229 For example, in programming courses, instructors and students may import
 230 the `davos` library into their notebooks to provide a simple means of ensur-
 231 ing their code will run on others’ machines. When combined with online
 232 notebook-based platforms like Google Colaboratory, `davos` provides a con-
 233 venient way to manage dependencies within a notebook, without requiring
 234 any software (beyond a web browser) to be installed on the students’ or in-
 235 structors’ systems. For the same reasons, `davos` also provides an elegant

236 means of sharing ready-to-run notebook-based demonstrations that install
237 their dependencies automatically.

238 Since its initial release, **davos** has found use in a variety of applications.
239 In addition to managing computing environments for multiple ongoing re-
240 search studies, **davos** is being used by both students and instructors in pro-
241 gramming and methods courses such as Storytelling with Data [18] (an open
242 course on data science, visualization, and communication) and Laboratory
243 in Psychological Science [19] (an open course on experimental and statistical
244 methods for psychology research) to simplify distributing lessons and sub-
245 mitting assignments, as well as in online demos such as **abstract2paper** [20]
246 (an example application of GPT-Neo [21, 22]) to share ready-to-run code
247 that installs dependencies automatically.

248 Our work also has several more subtle “advanced” use cases and poten-
249 tial impacts. Whereas Python’s built-in **import** statement is agnostic to
250 packages’ version numbers, **smuggle** statements (when combined with onion
251 comments) are version-sensitive. And because onion comments are parsed
252 at runtime, required package and their specified versions are installed in a
253 just-in-time manner. Thus, it is possible in most cases to **smuggle** a specific
254 package version or revision even if a different version has already been loaded.
255 This enables more complex uses that take advantage of multiple versions of
256 a package within a single interpreter session. This could be useful in cases
257 where specific features are added or removed from a package across differ-
258 ent versions, or in comparing the performance or functionality of particular
259 features across different versions of the same package.

260 A second advanced use case is in providing a proof-of-concept of how one
261 can add new “keyword-like” operators to the Python language by leverag-
262 ing notebooks’ error-handling mechanisms. This could lead to exciting new
263 tools that, like **davos**, extend the Python language in useful ways within
264 notebook-based environments. We note that our approach to adding the
265 **smuggle** keyword to Python when **davos** is imported into a notebook-based
266 environment also has the potential to be exploited for more nefarious pur-
267 poses. For example, a malicious user could use a similar approach (e.g.,
268 in a different library) to substantially change a notebook’s functionality by
269 adding new *unexpected* keyword-like objects (e.g., based around common ty-
270 pos). This could lead to difficult-to-predict changes in a notebook’s behavior
271 once the malicious library was imported. This highlights an important rea-
272 son why security-conscious users would be well-served to only make use of
273 libraries from trusted sources, or whose code is publicly available for review.

274 5. Conclusions

275 The `davos` library supports reproducible research by providing a novel
276 lightweight system for sharing notebook-based code. But perhaps the most
277 exciting uses of the `davos` library are those that we have *not* yet considered
278 or imagined. We hope that the Python community will find `davos` to pro-
279 vide a convenient means of managing project dependencies to facilitate code
280 sharing. We also hope that some of the more advanced applications of our
281 library might lead to new insights or discoveries.

282 Author Contributions

283 **Paxton C. Fitzpatrick:** Conceptualization, Methodology, Software,
284 Validation, Writing - Original Draft, Visualization. **Jeremy R. Manning:**
285 Conceptualization, Resources, Validation, Writing - Review & Editing, Su-
286 pervision, Funding acquisition.

287 Funding

288 Our work was supported in part by NSF grant number 2145172 to JRM.
289 The content is solely the responsibility of the authors and does not necessarily
290 represent the official views of our supporting organizations.

291 Declaration of Competing Interest

292 We wish to confirm that there are no known conflicts of interest associated
293 with this publication and there has been no significant financial support for
294 this work that could have influenced its outcome.

295 Acknowledgements

296 We acknowledge useful feedback and discussion from the students of
297 JRM's *Storytelling with Data* course (Winter, 2022 offering) who used pre-
298 liminary versions of our library in several assignments.

299 References

- 300 [1] G. van Rossum, Python reference manual, Department of Computer
301 Science [CS] (R 9525) (1995).
- 302 [2] Python Software Foundation, The Python Package Index (PyPI),
303 <https://pypi.org> (2003).

- 304 [3] conda-forge community, The conda-forge Project: Community-based
305 Software Distribution Built on the conda Package Format and Ecosys-
306 tem, <https://doi.org/10.5281/zenodo.4774217> (July 2015). doi:
307 [10.5281/zenodo.4774217](https://doi.org/10.5281/zenodo.4774217).
- 308 [4] N. Coghlan, D. Stufft, Version Identification and Dependency Specifica-
309 tion, PEP 440, Python Software Foundation (March 2013).
- 310 [5] B. Cannon, N. Smith, D. Stufft, Specifying minimum build system re-
311 quirements for python projects, PEP 518, Python Software Foundation
312 (May 2016).
- 313 [6] Anaconda, Inc., conda, <https://docs.conda.io> (2012).
- 314 [7] S. Eustace, Poetry: Python packaging and dependency management
315 made easy, <https://github.com/python-poetry/poetry> (December
316 2019).
- 317 [8] T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier,
318 J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov,
319 D. Avila, S. Abdalla, C. Willing, Jupyter Notebooks – a publish-
320 ing format for reproducible computational workflows, in: F. Loizides,
321 B. Schmidt (Eds.), Positioning and Power in Academic Publishing: Play-
322 ers, Agents and Agendas, IOS Press, Netherlands, 2016, pp. 87–90.
323 doi:[10.3233/978-1-61499-649-1-87](https://doi.org/10.3233/978-1-61499-649-1-87).
- 324 [9] R. P. Goldberg, Survey of virtual machine research, Computer 7 (6)
325 (1974) 34–45.
- 326 [10] Y. Altintas, C. Brecher, M. Weck, S. Witt, Virtual Machine Tool,
327 CIRP Annals 54 (2) (2005) 115–138. doi:[https://doi.org/10.1016/](https://doi.org/10.1016/S0007-8506(07)60022-5)
328 [S0007-8506\(07\)60022-5](https://doi.org/10.1016/S0007-8506(07)60022-5).
- 329 [11] M. Rosenblum, VMware’s Virtual Platform: A virtual machine monitor
330 for commodity PCs, in: IEEE Hot Chips Symposium, IEEE, 1999, pp.
331 185–196.
- 332 [12] D. Merkel, Docker: lightweight linux containers for consistent develop-
333 ment and deployment, Linux Journal 239 (2) (2014) 2.
- 334 [13] G. M. Kurtzer, V. Sochat, M. W. Bauer, Singularity: Scientific contain-
335 ers for mobility of compute, PLoS One 12 (5) (2017) e0177459.

- 336 [14] F. Pérez, B. E. Granger, IPython: a system for interactive scientific
337 computing, *Computing in science and engineering* 9 (3) (2007) 21–29.
338 [doi:10.1109/MCSE.2007.53](https://doi.org/10.1109/MCSE.2007.53).
- 339 [15] G. van Rossum, J. Lehtosalo, L. Langa, Type Hints, PEP 484, Python
340 Software Foundation (September 2014).
- 341 [16] L. Torvalds, J. Hamano, Git: Fast version control system, [https://](https://git.kernel.org/pub/scm/git/git.git)
342 git.kernel.org/pub/scm/git/git.git (April 2005).
- 343 [17] K. Thompson, Programming Techniques: Regular expression search al-
344 gorithm, *Communications of the ACM* 11 (6) (1968) 419–422. [doi:](https://doi.org/10.1145/363347.363387)
345 [10.1145/363347.363387](https://doi.org/10.1145/363347.363387).
- 346 [18] J. R. Manning, *Storytelling with Data*, [https://github.com/](https://github.com/ContextLab/storytelling-with-data)
347 [ContextLab/storytelling-with-data](https://github.com/ContextLab/storytelling-with-data) (June 2021). [doi:10.5281/](https://doi.org/10.5281/zenodo.5182775)
348 [zenodo.5182775](https://doi.org/10.5281/zenodo.5182775).
- 349 [19] J. Manning, *ContextLab/experimental-psychology: v1.0* (Spring, 2022),
350 [https://github.com/ContextLab/experimental-psychology/tree/](https://github.com/ContextLab/experimental-psychology/tree/v1.0)
351 [v1.0](https://github.com/ContextLab/experimental-psychology/tree/v1.0) (May 2022). [doi:10.5281/zenodo.6596762](https://doi.org/10.5281/zenodo.6596762).
- 352 [20] J. R. Manning, *abstract2paper*, [https://github.com/ContextLab/](https://github.com/ContextLab/abstract2paper)
353 [abstract2paper](https://github.com/ContextLab/abstract2paper) (June 2021).
- 354 [21] L. Gao, S. Biderman, S. Black, L. Golding, T. Hoppe, C. Foster,
355 J. Phang, H. He, A. Thite, N. Nabeshima, S. Presser, C. Leahy, *The*
356 *Pile: An 800GB Dataset of Diverse Text for Language Modeling*, arXiv
357 preprint arXiv:2101.00027 (2020).
- 358 [22] S. Black, L. Gao, P. Wang, C. Leahy, S. Biderman, *GPT-Neo: Large*
359 *Scale Autoregressive Language Modeling with Mesh-Tensorflow*, [http:](http://github.com/eleutherai/gpt-neo)
360 [//github.com/eleutherai/gpt-neo](http://github.com/eleutherai/gpt-neo) (2021).