

davos: a Python package “smuggler” for constructing lightweight reproducible notebooks

Paxton C. Fitzpatrick, Jeremy R. Manning*

*Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755*

Abstract

Reproducibility is a core requirement of modern scientific research. For computational research, reproducibility means that code should produce the same results, even when run on different systems. A standard approach to ensuring reproducibility entails packaging a project’s dependencies along with its primary code base. Existing solutions vary in how deeply these dependencies are specified, ranging from virtual environments, to containers, to virtual machines. Each of these existing solutions requires installing or setting up a system for running the desired code, increasing the complexity and time cost of sharing or engaging with reproducible science. Here, we propose a lighter-weight solution: the **davos** library. When used in combination with a notebook-based Python project, **davos** provides a mechanism for specifying (and automatically installing) the correct versions of the project’s dependencies. The **davos** library further ensures that those packages and specific versions are used every time the notebook’s code is executed. This enables researchers to share a complete reproducible copy of their code within a single Jupyter notebook file.

Keywords: Reproducibility, Open science, Python, Jupyter Notebook, Google Colaboratory, Package management

*Corresponding author

Email address: `Jeremy.R.Manning@Dartmouth.edu` (Jeremy R. Manning)

Metadata

Current code version

Nr.	Code metadata description	Metadata value
C1	Current code version	v0.1.1
C2	Permanent link to code/repository used for this code version	https://github.com/ContextLab/davos/tree/v0.1.1
C3	Code Ocean compute capsule	
C4	Legal Code License	MIT
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python, JavaScript, PyPI/pip, IPython, Jupyter, Ipykernel, PyZMQ. Additional tools used for tests: pytest, Selenium, Requests, mypy, GitHub Actions
C7	Compilation requirements, operating environments, and dependencies	Dependencies: Python ≥ 3.6 , packaging, setuptools. Supported OSes: MacOS, Linux, Unix-like. Supported IPython environments: Jupyter notebooks, JupyterLab, Google Colaboratory, Binder, IDE-based notebook editors.
C8	Link to developer documentation/manual	https://github.com/ContextLab/davos#readme
C9	Support email for questions	contextualdynamics@gmail.com

Table 1: Code metadata

1. Motivation and significance

The same computer code may not behave identically under different circumstances. For example, when code depends on external libraries, different versions of those libraries may function differently. Or when CPU or GPU instruction sets differ across machines, the same high-level code may be compiled into different machine instructions. Because executing identical code does not guarantee identical outcomes, code sharing alone is often insufficient for enabling researchers to reproduce each other’s work, or to collaborate on projects involving data collection or analysis.

Within the Python [1] community, external packages that are published in the most popular repositories [2, 3] are associated with version numbers and

12 tags that allow users to guarantee they are installing exactly the same code
13 across different computing environments [4]. While it is *possible* to manually
14 install the intended version of every dependency of a Python script or pack-
15 age, manually tracking down those dependencies can impose a substantial
16 burden on the user and create room for mistakes and inconsistencies. Fur-
17 ther, when dependency versions are left unspecified, replicating the original
18 computing environment becomes difficult or impossible.

19 Computational researchers and other programmers have developed a broad
20 set of approaches and tools to facilitate code sharing and reproducible out-
21 comes (Fig. 1). At one extreme, simply distributing a set of Python scripts
22 (`.py` files) may enable others to use or gain insights into the relevant work.
23 Because Python is installed by default on most modern operating systems,
24 for some projects, this may be sufficient. Another popular approach en-
25 tails creating Jupyter notebooks [8] that comprise a mix of text, executable
26 code, and embedded media. Notebooks may call or import external scripts
27 or libraries—or even intersperse snippets of other programming or markup
28 languages—in order to provide a more compact and readable experience for
29 users. Both of these systems (Python scripts and notebooks) provide a con-
30 venient means of sharing code, with the caveat that they do not specify the
31 computing environment in which the code is executed. Therefore the func-
32 tionality of code shared using these systems cannot be guaranteed across
33 different users or setups.

34 At another extreme, virtual machines [9, 10, 11] provide a hardware-
35 level simulation of the desired system. Virtual machines are typically iso-
36 lated, such that installing or running software on a virtual machine does
37 not impact the user’s primary operating system or computing environment.
38 Containers [e.g., 12, 13] provide a similar “isolated” experience. Although
39 containerized environments do not specify hardware-level operations, they
40 are typically packaged with a complete operating system, in addition to a
41 complete copy of Python and any relevant package dependencies. Virtual en-
42 vironments [e.g., 6, 7] also provide a computing environment that is largely
43 separated from the user’s main environment. They incorporate a copy of
44 Python and the target software’s dependencies, but virtual environments do
45 not specify or reproduce an operating system for the runtime environment.
46 Each of these systems (virtual machines, containers, and virtual environ-
47 ments) guarantees (to differing degrees—at the hardware level, operating
48 system level, and Python environment level, respectively) that the relevant
49 code will run similarly for different users. However, each of these systems
50 also relies on additional software that can be complex or resource-intensive
51 to install and use, creating potential barriers to both contributing to and
52 taking advantage of open science resources.

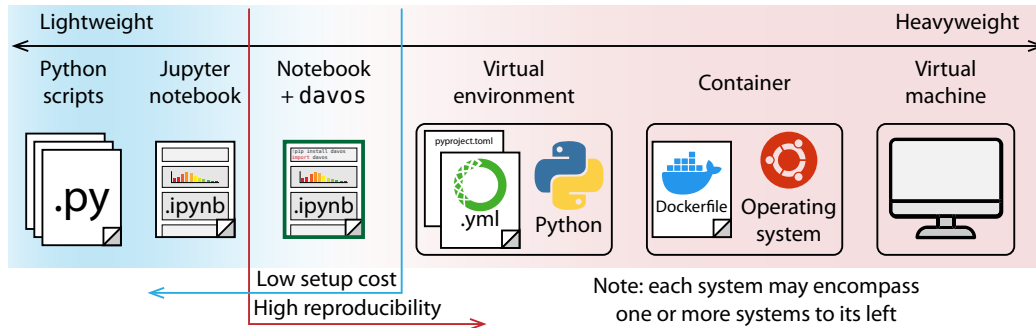


Figure 1: **Systems for sharing code within the Python ecosystem.** From left to right: plain-text **Python scripts** (`.py` files) provide the most basic “system” for sharing raw code. Scripts may reference external libraries, but those libraries must be manually installed on other users’ systems. Further, any checking needed to verify that the correct versions of those libraries were installed must also be performed manually. **Jupyter notebooks** (`.ipynb` files) comprise embedded text, executable code, and media (including rendered figures, code output, etc.). When the **davis** library is imported into a Jupyter notebook, the notebook’s functionality is extended to automatically install any required external libraries (at their correct versions, when specified). **Virtual environments** allow users to install an isolated copy of Python and all required dependencies. This typically entails distributing a configuration file (e.g., a `pyproject.toml` [5] or `environment.yml` file) that specifies all project dependencies (including version numbers of external libraries) alongside the primary code base. Users can then install a third-party tool [e.g., 6, 7] to read the file and build the environment. **Containers** provide a means of defining an isolated environment that includes a complete operating system (independent of the user’s operating system), in addition to (optionally) specifying a virtual environment or other configurations needed to provide the necessary computing environment. Containers are typically defined using specification files (e.g., a plain-text `Dockerfile`) that instruct the virtualization engine regarding how to build the containerized environment. **Virtual machines** provide a complete hardware-level simulation of the computing environment. In addition to simulating specific hardware, virtual machines (typically specified using binary image files) must also define operating system-level properties of the computing environment. Systems to the left of the blue vertical line entail sharing individual files, with no additional installation or configuration needed to run the target code. Systems to the right of the red vertical line support precise control over dependencies and versioning. Notebooks enhanced using the **davis** library are easily shareable and require minimal setup costs, while also facilitating high reproducibility by enabling precise control over project dependencies.

We designed **davos** to occupy a “sweet spot” between these extremes. **davos** is a notebook-installable package that adds functionality to the default notebook experience. Like standard Jupyter notebooks, **davos**-enhanced notebooks allow researchers to include text, executable code, and media within a single file. No further setup or installation is required, beyond what is needed to run standard Jupyter notebooks. And like virtual environments, **davos** provides a convenient mechanism for fully specifying (and installing, as needed) a complete set of Python dependencies, including package versions.

2. Software description

The **davos** package is named after Davos Seaworth, a smuggler referred to as “the Onion Knight” from the series *A Song of Ice and Fire* by George R. R. Martin [14]. The **smuggle** keyword-like operator implemented in **davos** is a play on Python’s **import** keyword: whereas importing can bring a package into the Python workspace within the existing rules and frameworks provided by the Python language, “smuggling” provides an alternative that expands the scope and reach of “importing.” Like the character Davos Seaworth (who became famous for smuggling onions through a blockade on his homeland), we use “onion” comments to precisely control how packages are smuggled into the Python workspace.

2.1. Software architecture

The **davos** package consists of two interdependent subpackages (see Fig. 2). The first, **davos.core**, comprises a set of modules that implement the bulk of the package’s core functionality, including pipelines for installing and validating packages, custom parsers for the **smuggle** statement (see Section 2.2.1) and onion comment (see Section 2.2.2), and a runtime interface for configuring **davos**’s behavior (see Section 2.2.3). However, certain critical aspects of this functionality require (often substantially) different implementations depending on properties of the notebook environment in which **davos** is used (e.g., whether the frontend is provided by Jupyter or Google Colaboratory, or which version of IPython [15] is used by the notebook kernel). To deal with this, environment-dependent parts of core features and behaviors are isolated and abstracted to “helper functions” in the **davos.implementations** subpackage. This second subpackage defines multiple, interchangeable versions of each helper function, organized into modules by the conditions that trigger their use. At runtime, **davos** detects various features in the notebook environment and selectively imports a single version of each helper function into the top-level **davos.implementations** namespace, allowing **davos.core** modules to access the proper implementations for the current

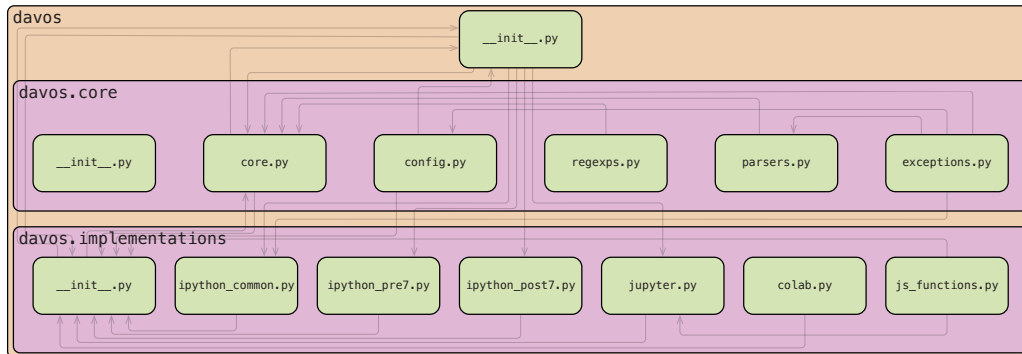


Figure 2: **Package structure.** The `davos` library comprises two interdependent subpackages. The `davos.core` subpackage includes modules for parsing `smuggle` statements and onion comments, installing and validating packages, and configuring `davos`’s behavior. The `davos.implementations` subpackage includes environment-specific modifications and features that are needed to support the core functionality across different notebook-based environments. Individual `.py` files are represented by lime rounded rectangles, and arrows denote dependencies (each arrow points to a file that import objects defined or specified at the arrow’s source).

notebook environment in a single, consistent location. An additional benefit of this design is that it allows maintainers, developers, and users to extend `davos` to support new, updated, or custom notebook variants by creating new `davos.implementations` modules that define their own versions of each helper function, modified from existing implementations as needed.

2.2. Software functionalities

2.2.1. The `smuggle` statement

Functionally, importing `davos` in an IPython notebook enables an additional Python keyword: “`smuggle`” (see Section 2.3 for details on how this works). The `smuggle` keyword-like object can be used as a drop-in replacement for Python’s built-in `import` keyword to load libraries, modules, and other objects into the current namespace. However, whereas `import` will fail if the requested package is not installed locally, `smuggle` statements can handle missing packages on the fly. If a smuggled package does not exist in the local environment, `davos` will install it automatically, expose its contents to Python’s `import` machinery, and load it into the namespace for immediate use.

108 2.2.2. The onion comment

109 For greater control over the behavior of **smuggle** statements, **davos** de-
110 fines an additional construct called the “onion comment.” An onion comment
111 is a special type of inline comment that may be placed on a line containing a
112 **smuggle** statement to customize how **davos** searches for the smuggled pack-
113 age locally and, if necessary, downloads and installs it. Onion comments
114 follow a simple format based on the “type comment” syntax introduced in
115 PEP 484 [16], and are designed to make managing packages with **davos** intu-
116 itive and familiar. To construct an onion comment, users provide the name
117 of the installer program (e.g., **pip**) and the same arguments one would use
118 to manually install the package as desired via the command line:

```
119 # enable smuggle statements
import davos

# if numpy is not installed locally, pip-install it and display verbose output
smuggle numpy as np # pip: numpy --verbose

# pip-install pandas without using or writing to the package cache
smuggle pandas as pd # pip: pandas --no-cache-dir

# install scipy from a relative local path, in editable mode
from scipy.stats smuggle ttest_ind # pip: -e ../../pkgs/scipy
```

120 Occasionally, a package’s distribution name (i.e., the name used when in-
121 stallng it) may differ from its top-level module name (i.e., the name used
122 when importing it). In such cases, an onion comment can be used to ensure
123 that **davos** installs the proper distribution if the smuggled package cannot
124 be found locally:

```
125 # package is named "python-dateutil" on PyPI, but imported as "dateutil"
smuggle dateutil # pip: python-dateutil

# package is named "scikit-learn" on PyPI, but imported as "sklearn"
from sklearn.decomposition smuggle PCA # pip: scikit-learn
```

126 Because onion comments may be constructed to specify any aspect of the
127 installer’s behavior, they provide a mechanism for precisely controlling how,
128 where, and when smuggled packages are installed. Critically, if an onion
129 comment includes a version specifier [4], **davos** will ensure that the version of
130 the package loaded into the notebook matches the specific version requested,
131 or satisfies the given version constraints. If the smuggled package exists
132 locally, **davos** will extract its version information from its metadata and
133 compare it to the specifier provided. If the two are incompatible (or no local
134 installation is found), **davos** will install and load a suitable version of the
135 package instead:

```

# specifically use matplotlib v3.4.2, pip-installing it if needed
smuggle matplotlib.pyplot as plt # pip: matplotlib==3.4.2

# use a version of seaborn no older than v0.9.1, but prior to v0.11
smuggle seaborn as sns # pip: seaborn>=0.9.1,<0.11

```

Onion comments can also be used to smuggle specific VCS references (e.g., Git [17] branches, commits, tags, etc.):

```

# use quail as the package existed on GitHub at commit 6c847a4
smuggle quail # pip: git+https://github.com/ContextLab/quail.git@6c847a4

```

davos processes onion comments internally before forwarding arguments to the installer program. In addition to preventing onion comments from being used as a vehicle for shell injection attacks, this design provides the user with additional control over which packages are imported. For example, each of the `-I/--ignore-installed`, `-U/--upgrade`, and `--force-reinstall` flags will cause **davos** to skip searching for a smuggled package locally before installing a new copy:

```

# install hypertools v0.7 without first checking for it locally
smuggle hypertools as hyp # pip: hypertools==0.7 --ignore-installed

# always install the latest version of requests, including pre-releases
from requests smuggle Session # pip: requests --upgrade --pre

```

Similarly, passing `--no-input` will temporarily enable **davos**'s non-interactive mode (see Section 2.2.3), and installing a smuggled package into a custom directory (`<dir>`) using the `--target <dir>` flag will cause **davos** to prepend `<dir>` to the module search path (i.e., `sys.path`), if necessary, so the package can be imported.

2.2.3. The *davos* config object

The **davos** config object provides a high-level interface for controlling various aspects of **davos**'s behavior. After importing **davos**, the `davos.config` object (a singleton) exposes configurable options as attributes that can be modified, checked at runtime, or displayed (see Sec. 3 for an illustrative example or Sec. 2.3 for implementation details and additional information). These include:

- **.active**: This attribute controls whether support for **smuggle** statements and onion comments) is enabled (**True**) or disabled (**False**). When **davos** is first imported, the **.active** attribute is set to **True**.

- 163 • `.auto_rerun`: This attribute controls how `davos` behaves when at-
164 tempting to `smuggle` a new version of a package that was previously
165 imported and cannot be reloaded. This can happen if the package in-
166 cludes extension modules that dynamically link C or C++ objects to
167 the Python interpreter itself, and if the code that generates those ob-
168 jects was changed between the previously imported and to-be-smuggled
169 versions. If this attribute is set to `True`, `davos` will automatically
170 restart the notebook kernel and rerun all code up to (and including)
171 the current `smuggle` statement. If `False` (the default), `davos` will in-
172 stead issue a warning, pause execution, and prompt the user to either
173 restart and rerun the notebook, or continue running with the previously
174 imported package version until the next time the kernel is restarted
175 manually. Note that, as of this writing, the `.auto_rerun` attribute is
176 not supported in Google Colaboratory notebooks.
 - 177 • `.confirm_install`: If `True` (default: `False`), `davos` will require user
178 confirmation before installing a smuggled package that do not yet exist
179 in the user's environment.
 - 180 • `.noninteractive`: Setting this attribute to `True` (default: `False`) en-
181 ables non-interactive mode, in which all user interactions (prompts and
182 dialogues) are disabled. Note that in non-interactive mode, the `con-`
183 `firm_install` option is set to `False`. If `auto_rerun` is `False` while
184 in non-interactive mode, `davos` will raise an exception if a smuggled
185 package cannot be reloaded, rather than prompting the user.
 - 186 • `.pip_executable`: This attribute's value specifies the path to the `pip`
187 executable used to install smuggled packages. The default is program-
188 matically determined from the Python environment and falls back to
189 `sys.executable -m pip` if no executable can be found.
 - 190 • `.suppress_stdout`: If this attribute is set to `True` (default: `False`),
191 `davos` suppresses printed (console) outputs from itself and the installer
192 program. This can be useful when smuggling packages that need to
193 install many dependencies and/or generate extensive output. However,
194 if the installer program throws an error, both the `stdout` and `stderr`
195 streams will still be displayed along with a stack trace.
- 196 The top-level `davos` namespace also defines convenience functions for setting
197 and checking whether `davos` is active (`davos.activate()`; `davos.deactivate()`;
198 `davos.is_active()`) as well as the `davos.configure()` function, which al-
199 lows setting multiple configuration options simultaneously.

200 2.3. Implementation details

201 Although **davos** is designed to appear to add a new Python keyword
202 to a notebook’s vocabulary, this illusion is actually a consequence of several
203 “hacks” that make use of the IPython backend for processing and executing
204 notebook code. Specifically, when **davos** is first imported, or when it acti-
205 vated after being set to an inactive state, two actions are triggered. First,
206 the **smuggle()** function is injected into the IPython user namespace. Second,
207 the **davos** parser is registered as a custom IPython input transformer.

208 IPython preprocesses all executed code as plain text before it is sent to
209 the Python compiler, in order to handle special constructs like **%magic** and
210 **!shell** commands. **davos** uses this process to transform **smuggle** statements
211 into syntactically valid Python code. The **davos** parser uses a regular expres-
212 sion to match lines of code containing **smuggle** statements (and, optionally,
213 onion comments), extract relevant information from their text, and replace
214 them with equivalent calls to the **smuggle()** function. For example, if a user
215 runs a notebook cell containing

```
216 smuggle numpy as np    # pip: numpy>1.16,<=1.20 -vv
```

217 the code that is actually executed by the Python interpreter would be

```
smuggle(name="numpy", as_="np", installer="pip",  
        args_str="\"numpy>1.16,<=1.20 -vv\"",  
        installer_kwargs={'editable': False,  
                          'spec': 'numpy>1.16,<=1.20',  
                          'verbosity': 2})
```

219 The call to the **smuggle()** function carries out **davos**’s central logic by de-
220 termining whether the smuggled package must be installed, carrying out the
221 installation if necessary, and subsequently loading it into the namespace.
222 This process is outlined in Figure 3. Because the **smuggle()** function is de-
223 fined in the notebook namespace, it is also possible (though never necessary)
224 to call it directly. Deactivating **davos** will delete the name “**smuggle**” from
225 the namespace, unless its value has been overwritten and no longer refers to
226 the **smuggle()** function. It will also deregister the **davos** parser from the
227 set of input transformers run when each notebook cell is executed. While
228 the overhead added by the **davos** parser is minimal, this may be useful, for
229 example, when optimizing or precisely profiling code.

230 3. Illustrative Example

231 Across different versions of a given package, functions may be changed,
232 depreciated, added, or renamed. In addition to changing the behaviors of

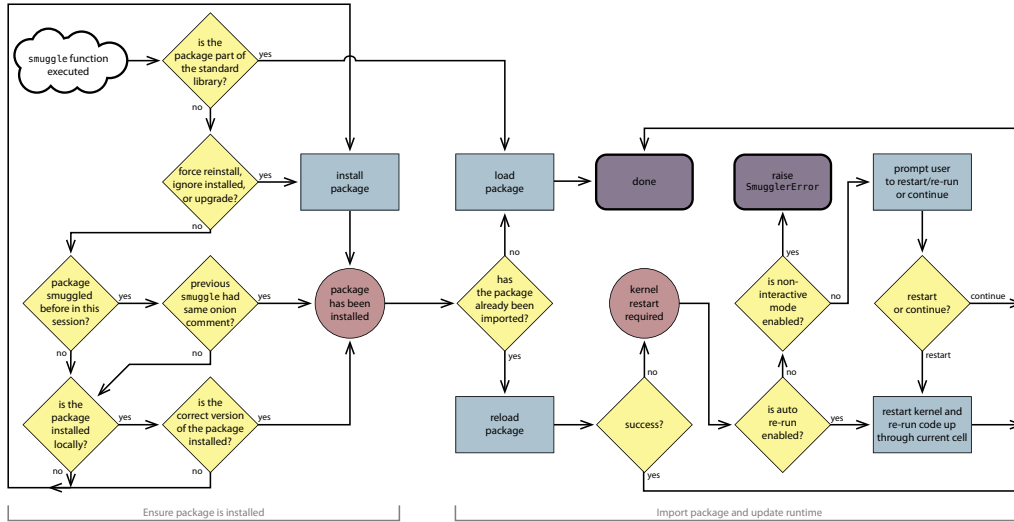


Figure 3: **smuggle() function algorithm.** At a high level, the `smuggle()` function may be conceptualized as following two basic steps. First (left), `davos` ensures that the correct version of the desired package has been installed, carrying out the installation automatically if needed. Second (right), `davos` imports the package and updates the current runtime environment.

active computations, saved objects created using one version of a package may be incompatible with another version of the same package. For example, the popular `pandas` [18] library prior to version 0.25.0 included the `Panel` datatype for storing 3-dimensional arrays. Since version 0.20.0, however, the `Panel` datatype has been deprecated, and the datatype was removed entirely in version 0.25.0. If one had saved a dataset in a `Panel` object (created using an older version of `pandas`), e.g., by serializing the `Panel` and saving it to disk using `pickle`, that dataset could not be read by any version of `pandas` from 0.25.0 or beyond. These incompatibilities are not limited solely to traditional forms of data. For example, saved model states and other objects may reference functions, attributes, classes, and other objects that are not present across all versions of their associated library.

The example provided in Figure 4 demonstrates how the `davos` library can be used to circumvent these incompatibilities by carefully controlling which versions of each package are used in different parts of the notebook. The example shows how a data and model saved using since-deprecated functionality in the `pandas` and `scikit-learn` [19] packages may be loaded in (using older versions of each package) and manipulated or analyzed (using more recent versions of each package).

After importing `davos` (line 1), the user first smuggles two utilities for

```

1  import davos
2
3  from os.path smuggle is_file
4  smuggle joblib                # pip: joblib<=1.2.0
5
6  davos.config.auto_rerun = True
7  smuggle numpy as np          # pip: numpy==1.21.6
8
9  if not is_file("~/datasets/data-new.csv"):
10     smuggle pandas as pd      # pip: pandas<0.25.0
11     tmp_data = pd.read_pickle("~/datasets/data-old.pkl")
12     tmp_data.to_frame().to_csv("~/datasets/data-new.csv")
13
14  smuggle pandas as pd         # pip: pandas==1.3.5
15
16  davos.configure(auto_rerun=False, suppress_stdout=True, noninteractive=True)
17  smuggle tensorflow as tf     # pip: tensorflow==2.9.2
18  from umap smuggle UMAP      # pip: umap-learn[plot,parametric_umap]==0.5.3
19  davos.configure(suppress_stdout=False, noninteractive=False)
20
21  smuggle matplotlib.pyplot as plt # pip: matplotlib==3.5.3
22  smuggle seaborn as sns       # pip: seaborn==0.12.1
23  smuggle quail                # pip: git+https://github.com/myfork/quail@6c847a4
24
25  davos.config.pip_executable = "~/envs/nb-server/bin/pip"
26  smuggle widgetsnbextension as _ # pip: widgetsnbextension==3.5.2
27  davos.config.pip_executable = "~/envs/nb-kernel/bin/pip"
28  smuggle ipywidgets           # pip: ipywidgets==7.6.5
29
30  from tqdm.notebook smuggle tqdm # pip: tqdm==4.62.3
31
32  data = pd.read_csv("~/datasets/data-new.csv", index_col=[0, 1])
33  smuggle sklearn              # pip: scikit-learn<0.22.0
34  transformer = joblib.load("~/models/text-transformer.joblib")
35  smuggle sklearn              # pip: scikit-learn==1.1.3

```

Figure 4: **Example use case for davos’s functionality.** Snippets from the example are also excerpted in the main text of Section 3.

253 interacting with local files in the code below. The `smuggle` statement in line
254 3 loads the `is_file()` function from the Python standard library's `os.path`
255 module. Standard library modules are included with all Python distributions,
256 so this line is functionally equivalent to an `import` statement and does not
257 need or benefit from an onion comment. Line 4 loads the `joblib` library [20],
258 installing it first, if necessary. Since `joblib`'s I/O interface has historically
259 remained stable and backwards-compatible across releases, requiring that
260 users have a particular exact version installed would likely be unnecessarily
261 restrictive. However, a *future* release might introduce some breaking change.
262 The onion comment in line 4 helps ensure the analysis notebook continues
263 to run properly in the future, by limiting allowable versions to those already
264 released when the code was written:

```
1  import davos
2
3  from os.path smuggle is_file
265 4  smuggle joblib                # pip: joblib<=1.2.0
```

266 Line 6 then uses the `davos.config` object to enable `davos`'s `auto_rerun`
267 option before smuggling the next two packages: `NumPy` [21] and `pandas`. Be-
268 cause these libraries rely heavily on custom C data types, loading the partic-
269 ular versions from the onion comments may require restarting the notebook
270 kernel if different versions had been previously imported during the same
271 interpreter session (see Section 2.2.3).

```
6  davos.config.auto_rerun = True
272 7  smuggle numpy as np          # pip: numpy==1.21.6
```

273 Setting the `auto_rerun` attribute to `True` also benefits the (potential) instal-
274 lation of `pandas` on the next lines:

```
9  if not is_file("~/datasets/data-new.csv"):
10     smuggle pandas as pd          # pip: pandas<0.25.0
11     tmp_data = pd.read_pickle("~/datasets/data-old.pkl")
12     tmp_data.to_frame().to_csv("~/datasets/data-new.csv")
13
275 14  smuggle pandas as pd          # pip: pandas==1.3.5
```

276 If we suppose that the data contained in `data-old.pkl` is stored in a `Panel`
277 object (i.e., created in a version of `pandas` prior to 0.25.0), then we would
278 not be able to load in that object with newer versions of `pandas`. Line 10
279 ensures that an older version of `pandas` will be imported, enabling the data
280 to be read in (and saved out to a `.csv` file, which is compatible with newer
281 versions of `pandas`).

282 Newer versions of `pandas` have brought substantial performance improve-
283 ments, added new functionality, fixed bugs, etc. Although the original dataset

284 had to be read in using an older version of the package, we can take advantage
285 of those newer changes using a new `smuggle` statement and onion comment
286 on line 14 (which specifies that version 1.3.5 should be imported). Since
287 `pandas` had already been imported into the current workspace (on line 10),
288 the runtime must be restarted in order to gain access to the newer version of
289 `pandas`. The `.auto_rerun` flag set on line 6 enables this process to happen
290 automatically, and saving out the dataset to a `.csv` file in lines 9–12 ensures
291 that the older version of `pandas` does not need to be reinstalled.

292 Next, line 16 uses the `davos.configure()` function to disable the `auto_rerun`
293 option and simultaneously enable two other options: `suppress_stdout`
294 and `noninteractive`. With these options enabled, lines 17–18 `smuggle`
295 `TensorFlow` [22], a powerful end-to-end platform for building and working
296 with machine learning models, and `UMAP` [23], a library that implements a
297 family of related manifold learning techniques. The onion comment in line
298 18 also specifies that `UMAP` should be installed with the optional requirements
299 needed for its “plot” and “parametric_umap” features. Together, these two
300 packages depend on 36 other unique packages, most of which have dependen-
301 cies of their own. And if many of these are not already installed in the user’s
302 environment, lines 17–18 could take several minutes to run. Enabling the
303 `noninteractive` option ensures that the installation will continue automat-
304 ically without user input during that time. Enabling `suppress_stdout` also
305 suppresses console outputs from the installer that might otherwise distract
306 from other more important outputs.

```
16  davos.configure(auto_rerun=False, suppress_stdout=True, noninteractive=True)
17  smuggle tensorflow as tf           # pip: tensorflow==2.9.2
18  from umap smuggle UMAP            # pip: umap-learn[plot,parametric_umap]==0.5.3
```

308 After re-enabling these two options (line 19), the user next smuggles spe-
309 cific versions of three plotting libraries: `Matplotlib` [24], `seaborn` [25], and
310 `Quail` [26] (lines 21–23). Because the first two are requirements of `UMAP`’s op-
311 tional “plot” feature, they will have already been installed by line 18, though
312 possibly as different versions than those specified in the onion comments on
313 lines 21 and 22. If the installed and specified versions are the same, these
314 `smuggle` statements will function like standard `import` statements to load
315 the libraries into the notebook namespace. If they differ, `davos` will down-
316 load the requested versions in place of the installed versions before doing
317 so.

318 Line 23 uses an onion comment to specify that `Quail` should be installed
319 directly from a specific GitHub commit (6c847a4). This ability to install
320 packages directly from GitHub repositories can enable developers to use
321 forked or modified versions of other libraries in their notebooks, even if those
322 versions have not been officially released.

```

19  davos.configure(suppress_stdout=False, noninteractive=False)
20
21  smuggle matplotlib.pyplot as plt      # pip: matplotlib==3.5.3
22  smuggle seaborn as sns                # pip: seaborn==0.12.1
323 23  smuggle quail                        # pip: git+https://github.com/myfork/quail@6c847a4

```

324 In lines 25–28, we demonstrate another aspect of `davos`’s functionality
 325 that supports more advanced installation scenarios. The `ipywidgets` [27]
 326 package provides an API for creating these widgets with Python code, and
 327 the `widgetsnbextension` package provides the machinery for the notebook
 328 frontend to display them.

```

25  davos.config.pip_executable = "~/envs/nb-server/bin/pip"
26  smuggle widgetsnbextension as _      # pip: widgetsnbextension==3.5.2
27  davos.config.pip_executable = "~/envs/nb-kernel/bin/pip"
28  smuggle ipywidgets                  # pip: ipywidgets==7.6.5
29
329 30  from tqdm.notebook import tqdm      # pip: tqdm==4.62.3

```

330 A complication is that `ipywidgets` must be installed in the same environment
 331 as the IPython kernel, whereas `widgetsnbextension` must be installed in the
 332 environment that houses the Jupyter notebook server. In standard setups,
 333 these two environments are the same. However, a common “advanced” ap-
 334 proach is to run the notebook server from a base environment, with additional
 335 environments each providing their own separate, interchangeable IPython
 336 kernels. To accomodate this multi-environment scenario, on lines 25 and 27
 337 we control which environments each package should be installed to. Once
 338 these two packages are installed and imported, line 30 smuggles `tqdm` [28],
 339 which display progress bars to provide status updates for running code. In
 340 Jupyter notebooks, the `tqdm.notebook` module can be imported to enable
 341 aesthetically pleasing progress bars that are displayed via `ipywidgets`, if
 342 that package is installed and importable. Therefore, to take advantage of
 343 this feature, it was important to smuggle `tqdm` after ensuring the `ipywid-`
 344 `gets` package was available.

345 Next, we load in the reformatted dataset (line 32) and pre-trained model
 346 (line 34) that we wish to use in our analysis. In our hypothetical example,
 347 we can suppose that the model was provided as a `scikit-learn` `Pipeline`
 348 object that passes data through two pretrained models in succession. First,
 349 a trained `CountVectorizer` instance converts text data to an array of word
 350 counts. Second, the word counts are passed to a topic model [29] using a
 351 pretrained `LatentDirichletAllocation` instance.

352 Let us suppose that the `Pipeline` object had been saved by its original
 353 creator using the `joblib` library, as `scikit-learn`’s documentation recom-
 354 mends. Because `joblib` uses the `pickle` protocol internally, the ability to

355 save and load pre-trained models is not guaranteed across different `scikit-`
356 `learn` versions. For example, suppose that the `Pipeline` object was created
357 using `scikit-learn` v0.21.3). In that version of `scikit-learn`, the `Latent-`
358 `DirichletAllocation` class is defined in `sklearn.decomposition.online_-`
359 `lda`. However, in version 0.22.0, that module was renamed to `_online_lda`,
360 and in versions 0.22.1 and higher, the name has been reverted back to `_lda`.

361 In order to correctly load the model that includes the pretrained `Latent-`
362 `DirichletAllocation` instance, in line 33, the user first smuggles a version
363 of `scikit-learn` prior to v0.22.0 (i.e., before the first time the relevant
364 module's name was changed). Once the model is loaded and reconstructed
365 in memory from a compatible package version (line 34), we upgrade to a
366 newer version of `scikit-learn` in line 35. Taken together, the code in Fig-
367 ure 4 shows how `davos` can enable users to load in data and models that
368 are incompatible with newer versions of `pandas` and `scikit-learn`, but still
369 *analyze* and the data and model output using the latest approaches and
370 implementations.

371 4. Impact

372 Like virtual environments, containers, and virtual machines, the `davos`
373 library (when used in conjunction with Jupyter notebooks) provides a light-
374 weight mechanism for sharing code and ensuring reproducibility across users
375 and computing environments (Fig. 1). Further, `davos` enables users to fully
376 specify (and install, as needed) any project dependencies within the same
377 notebook. This provides a system whereby executable code (along with text
378 and media) *and* code for setting up and configuring the project dependencies,
379 may be combined within a single notebook file.

380 We designed `davos` for use in research applications. For example, in many
381 settings, `davos` may be used as a drop-in replacement for more-difficult-to-
382 set-up virtual environments, containers, and/or virtual machines. For re-
383 searchers, this lowers barriers to both sharing code. By eliminating most
384 of the setup costs of reconstructing the original researchers' computing en-
385 vironment, `davos` also lowers barriers to entry for members of the scientific
386 community and the public who seek to benefit from shared code.

387 Beyond research applications, `davos` is also useful in pedagogical set-
388 tings. For example, in programming courses, instructors and students may
389 use the `davos` library to ensure their notebooks will run correctly on oth-
390 ers' machines. When combined with online notebook-based platforms like
391 Google Colaboratory, `davos` provides a convenient way to manage depen-
392 dencies within a notebook, without requiring any software (beyond a web
393 browser) to be installed on the students' or instructors' systems. For the

394 same reasons, **davos** also provides an elegant means of sharing ready-to-run
395 notebook-based demonstrations or tutorials that install their dependencies
396 automatically.

397 Since its initial release, **davos** has found use in a variety of applications.
398 In addition to managing computing environments for multiple ongoing re-
399 search studies, **davos** is being used by both students and instructors in pro-
400 gramming and methods courses such as Storytelling with Data [30] (an open
401 course on data science, visualization, and communication) and Laboratory
402 in Psychological Science [31] (an open course on experimental and statistical
403 methods for psychology research) to simplify distributing lessons and sub-
404 mitting assignments, as well as in online demos such as **abstract2paper** [32]
405 (an example application of GPT-Neo [33, 34]) to share ready-to-run code
406 that installs dependencies automatically.

407 Our work also has several more subtle “advanced” use cases and potential
408 impacts. Whereas Python’s built-in **import** statement is agnostic to pack-
409 ages’ version information, **smuggle** statements (when combined with onion
410 comments) are version-sensitive. And because onion comments are parsed
411 at runtime, required packages and their specified versions are installed in a
412 just-in-time manner. Thus, it is possible in most cases to **smuggle** a specific
413 package version or revision even if a different version has already been loaded.
414 This enables more complex uses that take advantage of multiple versions of a
415 package within a single interpreter session (e.g., see Section 3). This could be
416 useful in cases where specific features are added or removed from a package
417 across different versions, or in comparing the performance or functionality of
418 particular features across different versions of the same package.

419 A second more subtle impact of our work is in providing a proof-of-concept
420 of how the ability to add new “keyword-like” operators to the Python lan-
421 guage could be specifically useful to researchers. With **davos**, we accomplish
422 this by leveraging IPython notebooks’ internal code parsing and execution
423 machinery. We note that, while other popular packages similarly use these
424 mechanisms to providing notebook-specific functionality (e.g., [24, 35]), this
425 approach also has the potential to be exploited for more nefarious purposes.
426 For example, a malicious user could design a Python library that, when
427 imported, substantially changes the notebook’s functionality by adding new
428 *unexpected* keyword-like objects (e.g., based around common typos). We also
429 note that this implementation approach means **davos**’s functionality is cur-
430 rently restricted to IPython notebook environments, as would be that of any
431 potential similar tools that enable user-defined keywords. However, there has
432 been early-stage discussion of providing this sort of syntactic customizability
433 as a core feature of the Python language, including a draft proposal [36]. In
434 addition to enabling **davos** to be extended for use outside of notebooks, this

435 could lead to exciting new tools that, like **davos**, extend the Python language
436 in useful and more secure ways.

437 **5. Conclusions**

438 The **davos** library supports reproducible research by providing a novel
439 lightweight system for sharing notebook-based code. But perhaps the most
440 exciting uses of the **davos** library are those that we have *not* yet considered
441 or imagined. We hope that the Python community will find **davos** to pro-
442 vide a convenient means of managing project dependencies to facilitate code
443 sharing. We also hope that some of the more advanced applications of our
444 library might lead to new insights or discoveries.

445 **Author Contributions**

446 **Paxton C. Fitzpatrick:** Conceptualization, Methodology, Software,
447 Validation, Writing - Original Draft, Visualization. **Jeremy R. Manning:**
448 Conceptualization, Resources, Validation, Writing - Review & Editing, Su-
449 pervision, Funding acquisition.

450 **Funding**

451 Our work was supported in part by NSF grant number 2145172 to JRM.
452 The content is solely the responsibility of the authors and does not necessarily
453 represent the official views of our supporting organizations.

454 **Declaration of Competing Interest**

455 We wish to confirm that there are no known conflicts of interest associated
456 with this publication and there has been no significant financial support for
457 this work that could have influenced its outcome.

458 **Acknowledgements**

459 We acknowledge useful feedback and discussion from the students of
460 JRM's *Storytelling with Data* course (Winter, 2022 offering) who used pre-
461 liminary versions of our library in several assignments.

462 References

- 463 [1] G. van Rossum, Python reference manual, Department of Computer
464 Science [CS] (R 9525) (1995).
- 465 [2] Python Software Foundation, The Python Package Index (PyPI),
466 <https://pypi.org> (2003).
- 467 [3] conda-forge community, The conda-forge Project: Community-based
468 Software Distribution Built on the conda Package Format and Ecosys-
469 tem, <https://doi.org/10.5281/zenodo.4774217> (July 2015). doi:
470 [10.5281/zenodo.4774217](https://doi.org/10.5281/zenodo.4774217).
- 471 [4] N. Coghlan, D. Stufft, Version Identification and Dependency Specifica-
472 tion, PEP 440, Python Software Foundation (March 2013).
- 473 [5] B. Cannon, N. Smith, D. Stufft, Specifying Minimum Build System Re-
474 quirements for Python Projects, PEP 518, Python Software Foundation
475 (May 2016).
- 476 [6] Anaconda, Inc., conda, <https://docs.conda.io> (2012).
- 477 [7] S. Eustace, Poetry: Python packaging and dependency management
478 made easy, <https://github.com/python-poetry/poetry> (December
479 2019).
- 480 [8] T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier,
481 J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov,
482 D. Avila, S. Abdalla, C. Willing, Jupyter Notebooks – a publish-
483 ing format for reproducible computational workflows, in: F. Loizides,
484 B. Schmidt (Eds.), Positioning and Power in Academic Publishing: Play-
485 ers, Agents and Agendas, IOS Press, Netherlands, 2016, pp. 87–90.
486 doi:[10.3233/978-1-61499-649-1-87](https://doi.org/10.3233/978-1-61499-649-1-87).
- 487 [9] R. P. Goldberg, Survey of virtual machine research, Computer 7 (6)
488 (1974) 34–45.
- 489 [10] Y. Altintas, C. Brecher, M. Weck, S. Witt, Virtual Machine Tool,
490 CIRP Annals 54 (2) (2005) 115–138. doi:[https://doi.org/10.1016/](https://doi.org/10.1016/S0007-8506(07)60022-5)
491 [S0007-8506\(07\)60022-5](https://doi.org/10.1016/S0007-8506(07)60022-5).
- 492 [11] M. Rosenblum, VMware’s Virtual Platform: A virtual machine monitor
493 for commodity PCs, in: IEEE Hot Chips Symposium, IEEE, 1999, pp.
494 185–196.

- [12] D. Merkel, Docker: lightweight linux containers for consistent development and deployment, *Linux Journal* 239 (2) (2014) 2.
- [13] G. M. Kurtzer, V. Sochat, M. W. Bauer, Singularity: Scientific containers for mobility of compute, *PLoS One* 12 (5) (2017) e0177459.
- [14] G. R. R. Martin, *A song of fire and ice*, HarperVoyager, London, England, 2012.
- [15] F. Pérez, B. E. Granger, IPython: a system for interactive scientific computing, *Computing in science and engineering* 9 (3) (2007) 21–29. doi:10.1109/MCSE.2007.53.
- [16] G. van Rossum, J. Lehtosalo, L. Langa, Type Hints, PEP 484, Python Software Foundation (September 2014).
- [17] L. Torvalds, J. Hamano, Git: Fast version control system, <https://git.kernel.org/pub/scm/git/git.git> (April 2005).
- [18] W. McKinney, Data Structures for Statistical Computing in Python, in: S. van der Walt, J. Millman (Eds.), *Proceedings of the 9th Python in Science Conference*, 2010, pp. 56–61. doi:10.25080/Majors-92bf1922-00a.
- [19] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: machine learning in Python, *Journal of Machine Learning Research* 12 (2011) 2825–2830.
- [20] G. Varoquaux, Joblib: Computing with Python functions, <https://github.com/joblib/joblib> (July 2010).
- [21] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. F. del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, T. E. Oliphant, Array programming with NumPy, *Nature* 585 (7825) (2020) 357–362. doi:10.1038/s41586-020-2649-2.
- [22] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow,

- 528 A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kud-
529 lur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah,
530 M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker,
531 V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wat-
532 tenberg, M. Wicke, Y. Yu, X. Zheng, [TensorFlow: Large-Scale Ma-
533 chine Learning on Heterogeneous Systems](#), software available from ten-
534 sorflow.org (2015).
535 URL <https://www.tensorflow.org/>
- 536 [23] L. McInnes, J. Healy, N. Saul, L. Großberger, UMAP: Uniform Manifold
537 Approximation and Projection, *Journal of Open Source Software* 3 (29)
538 (2018) 861. doi:<https://doi.org/10.21105/joss.00861>.
- 539 [24] J. D. Hunter, Matplotlib: A 2D graphics environment, *Computing in
540 Science and Engineering* 9 (3) (2007) 90–95. doi:[10.1109/MCSE.2007.
541 55](https://doi.org/10.1109/MCSE.2007.55).
- 542 [25] M. L. Waskom, seaborn: statistical data visualization, *Journal of Open
543 Source Software* 6 (60) (2021) 3021. doi:[10.21105/joss.03021](https://doi.org/10.21105/joss.03021).
- 544 [26] A. C. Heusser, P. C. Fitzpatrick, C. E. Field, K. Ziman, J. R. Manning,
545 Quail: a Python toolbox for analyzing and plotting free recall data,
546 *Journal of Open Source Software* 10.21105/joss.00424 (2017).
- 547 [27] J. Frederic, J. Grout, Jupyter Widgets Contributors, ipywidgets: In-
548 teractive Widgets for the Jupyter Notebook, [https://github.com/
549 jupyter-widgets/ipywidgets](https://github.com/jupyter-widgets/ipywidgets) (August 2015).
- 550 [28] C. da Costa-Luis, S. K. Larroque, K. Altendorf, H. Mary, richardsheri-
551 dan, M. Korobov, N. Raphael, I. Ivanov, M. Bargull, N. Rodrigues,
552 G. Chen, A. Lee, C. Newey, CrazyPython, JC, M. Zugnoni, M. D.
553 Pagel, mjstevens777, M. Dektyarev, A. Rothberg, A. Plavin, D. Pan-
554 teleit, F. Dill, FichteFoll, G. Sturm, HeoHeo, H. van Kemenade, J. Mc-
555 Cracken, MapleCCC, M. Nordlund, tqdm: A Fast, Extensible Progress
556 Bar for Python and CLI, <https://github.com/tqdm/tqdm> (September
557 2022). doi:[10.5281/zenodo.595120](https://doi.org/10.5281/zenodo.595120).
- 558 [29] D. M. Blei, A. Y. Ng, M. I. Jordan, Latent dirichlet allocation, *Journal
559 of Machine Learning Research* 3 (2003) 993–1022.
- 560 [30] J. R. Manning, *Storytelling with Data*, [https://github.com/
561 ContextLab/storytelling-with-data](https://github.com/ContextLab/storytelling-with-data) (June 2021). doi:[10.5281/
562 zenodo.5182775](https://doi.org/10.5281/zenodo.5182775).

- 563 [31] J. Manning, ContextLab/experimental-psychology: v1.0 (Spring, 2022),
 564 [https://github.com/ContextLab/experimental-psychology/tree/](https://github.com/ContextLab/experimental-psychology/tree/v1.0)
 565 [v1.0](https://github.com/ContextLab/experimental-psychology/tree/v1.0) (May 2022). doi:10.5281/zenodo.6596762.
- 566 [32] J. R. Manning, abstract2paper, [https://github.com/ContextLab/](https://github.com/ContextLab/abstract2paper)
 567 [abstract2paper](https://github.com/ContextLab/abstract2paper) (June 2021). doi:10.5281/zenodo.7261831.
- 568 [33] L. Gao, S. Biderman, S. Black, L. Golding, T. Hoppe, C. Foster,
 569 J. Phang, H. He, A. Thite, N. Nabeshima, S. Presser, C. Leahy, The
 570 Pile: An 800GB Dataset of Diverse Text for Language Modeling, arXiv
 571 preprint arXiv:2101.00027 (2020).
- 572 [34] S. Black, L. Gao, P. Wang, C. Leahy, S. Biderman, GPT-Neo: Large
 573 Scale Autoregressive Language Modeling with Mesh-Tensorflow, [http:](http://github.com/eleutherai/gpt-neo)
 574 [//github.com/eleutherai/gpt-neo](http://github.com/eleutherai/gpt-neo) (2021).
- 575 [35] A. C. Heusser, K. Ziman, L. L. W. Owen, J. R. Manning, HyperTools:
 576 a Python toolbox for gaining geometric insights into high-dimensional
 577 data, Journal of Machine Learning Research 18 (152) (2018) 1–6.
- 578 [36] M. Shannon, Syntactic Macros, Draft PEP 638, Python Software Foun-
 579 dation (September 2020).