

davos: a Python package “smuggler” for constructing lightweight reproducible notebooks

Paxton C. Fitzpatrick, Jeremy R. Manning*

*Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755*

Abstract

Reproducibility is a core requirement of modern scientific research. For computational research, reproducibility means that code should produce the same results, even when run on different systems. A standard approach to ensuring reproducibility entails packaging a project’s dependencies along with its primary code base. Existing solutions vary in how deeply these dependencies are specified, ranging from virtual environments, to containers, to virtual machines. Each of these existing solutions requires installing or setting up a system for running the desired code, increasing the complexity and time cost of sharing or engaging with reproducible science. Here, we propose a lighter-weight solution: the **davos** package. When used in combination with a notebook-based Python project, **davos** provides a mechanism for specifying (and automatically installing) the correct versions of the project’s dependencies. The **davos** package further ensures that those packages and specific versions are used every time the notebook’s code is executed. This enables researchers to share a complete reproducible copy of their code within a single Jupyter notebook file.

Keywords: Reproducibility, Open science, Python, Jupyter Notebook, Google Colaboratory, Package management

*Corresponding author

Email address: `Jeremy.R.Manning@Dartmouth.edu` (Jeremy R. Manning)

Metadata

Current code version

Nr.	Code metadata description	Metadata value
C1	Current code version	v0.1.1
C2	Permanent link to code/repository used for this code version	https://github.com/ContextLab/davos/tree/v0.1.1
C3	Code Ocean compute capsule	
C4	Legal Code License	MIT
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python, JavaScript, PyPI/pip, IPython, Jupyter, Ipykernel, PyZMQ. Additional tools used for tests: pytest, Selenium, Requests, mypy, GitHub Actions
C7	Compilation requirements, operating environments, and dependencies	Dependencies: Python ≥ 3.6 , packaging, setuptools. Supported OSes: MacOS, Linux, Unix-like. Supported IPython environments: Jupyter notebooks, JupyterLab, Google Colaboratory, Binder, IDE-based notebook editors.
C8	Link to developer documentation/manual	https://github.com/ContextLab/davos#readme
C9	Support email for questions	contextualdynamics@gmail.com

Table 1: Code metadata

1. Motivation and significance

The same computer code may not behave identically under different circumstances. For example, when code depends on external packages, different versions of those packages may function differently. Or when CPU or GPU instruction sets differ across machines, the same high-level code may be compiled into different machine instructions. Because executing identical code does not guarantee identical outcomes, code sharing alone is often insufficient for enabling researchers to reproduce each other’s work, or to collaborate on projects involving data collection or analysis.

Within the Python [1] community, external packages that are published in the most popular repositories [2, 3] are associated with version numbers and tags that allow users to guarantee they are installing exactly the same code across different computing environments [4]. While it is *possible* to manually install the intended version of every dependency of a Python script or package, manually tracking down those dependencies can impose a substantial burden on the user and create room for mistakes and inconsistencies. Further, when dependency versions are left unspecified, replicating the original computing environment becomes difficult or impossible.

Computational researchers and other programmers have developed a broad set of approaches and tools to facilitate code sharing and reproducible outcomes (Fig. 1). At one extreme, simply distributing a set of Python scripts (.py files) may enable others



Figure 1: **Systems for sharing code within the Python ecosystem.** From left to right: plain-text **Python scripts** (`.py` files) provide the most basic “system” for sharing raw code. Scripts may reference external packages, but those packages must be manually installed on other users’ systems. Further, any checking needed to verify that the correct versions of those packages were installed must also be performed manually. **Jupyter notebooks** (`.ipynb` files) comprise embedded text, executable code, and media (including rendered figures, code output, etc.). When the **dapos** package is imported into a Jupyter notebook, the notebook’s functionality is extended to automatically install any required external packages (at their correct versions, when specified). **Virtual environments** allow users to install an isolated copy of Python and all required dependencies. This typically entails distributing a configuration file (e.g., a `pyproject.toml` [5] or `environment.yaml` file) that specifies all project dependencies (including version numbers of external packages) alongside the primary code base. Users can then install a third-party tool [e.g., 6, 7] to read the file and build the environment. **Containers** provide a means of defining an isolated environment that includes a complete operating system (independent of the user’s operating system), in addition to (optionally) specifying a virtual environment or other configurations needed to provide the necessary computing environment. Containers are typically defined using specification files (e.g., a plain-text `Dockerfile`) that instruct the virtualization engine regarding how to build the containerized environment. **Virtual machines** provide a complete hardware-level simulation of the computing environment. In addition to simulating specific hardware, virtual machines (typically specified using binary image files) must also define operating system-level properties of the computing environment. Systems to the left of the blue vertical line entail sharing individual files, with no additional installation or configuration needed to run the target code. Systems to the right of the red vertical line support precise control over dependencies and versioning. Notebooks enhanced using the **dapos** package are easily shareable and require minimal setup costs, while also facilitating high reproducibility by enabling precise control over project dependencies.

to use or gain insights into the relevant work. Because Python is installed by default on most modern operating systems, for some projects, this may be sufficient. Another popular approach entails creating Jupyter notebooks [8] that comprise a mix of text, executable code, and embedded media. Notebooks may call or import external scripts or packages—or even intersperse snippets of other programming or markup languages—in order to provide a more compact and readable experience for users. Both of these systems (Python scripts and notebooks) provide a convenient means of sharing code, with the caveat that they do not specify the computing environment in which the code is executed. Therefore the functionality of code shared using these systems cannot be guaranteed across different users or setups.

At another extreme, virtual machines [9, 10, 11] provide a hardware-level simulation of the desired system. Virtual machines are typically isolated, such that installing or running software on a virtual machine does not impact the user’s primary operating system or computing environment. Containers [e.g., 12, 13] provide a similar “isolated” experience. Although containerized environments do not specify hardware-level operations, they are typically packaged with a complete operating system, in addition to a complete copy of Python and any relevant package dependencies. Virtual environments [e.g., 6, 7] also provide a computing environment that is largely separated from the user’s main environment. They incorporate a copy of Python and the target software’s dependencies, but virtual environments do not specify or reproduce an operating system for the runtime environment. Each of these systems (virtual machines, containers, and virtual environments) guarantees (to differing degrees—at the hardware level, operating system level, and Python environment level, respectively) that the relevant code will run similarly for different users. However, each of these systems also relies on additional software that can be complex or resource-intensive to install and use, creating potential barriers to both contributing to and taking advantage of open science resources.

We designed **davos** to occupy a “sweet spot” between these extremes. **davos** is a notebook-installable package that adds functionality to the default notebook experience. Like standard Jupyter notebooks, **davos**-enhanced notebooks allow researchers to include text, executable code, and media within a single file. No further setup or installation is required, beyond what is needed to run standard Jupyter notebooks. And like virtual environments, **davos** provides a convenient mechanism for fully specifying (and installing, as needed) a complete set of Python dependencies, including specific package versions.

2. Software description

The **davos** package is named after Davos Seaworth, a smuggler referred to as “the Onion Knight” from the series *A Song of Ice and Fire* by George R. R. Martin [14]. The **smuggle** keyword provided by **davos** is a play on Python’s **import** keyword: whereas importing can load a package into the Python workspace within the existing rules and frameworks provided by the Python language, “smuggling” provides an alternative that expands the scope and reach of “importing.” Like the character Davos Seaworth (who became famous for smuggling onions through a blockade on his homeland), we use “onion” comments to precisely control how packages are smuggled into the Python workspace.

2.1. Software architecture

The **davos** package consists of two interdependent subpackages (see Fig. 2). The first, **davos.core**, comprises a set of modules that implement the bulk of the package’s core




Figure 2: **Package structure.** The `davos` package comprises two interdependent subpackages. The `davos.core` subpackage includes modules for parsing `smuggle` statements and onion comments, installing and validating packages, and configuring `davos`’s behavior. The `davos.implementations` subpackage includes environment-specific modifications and features that are needed to support the core functionality across different notebook-based environments. Individual modules (i.e., `.py` files) are represented by lime rounded rectangles, and arrows denote dependencies (each arrow points to a module that imports objects defined in the module at the arrow’s source).

functionality, including pipelines for installing and validating packages, custom parsers for the `smuggle` statement (see Sec. 2.2.1) and onion comment (see Sec. 2.2.2), and a runtime interface for configuring `davos`’s behavior (see Sec. 2.2.3). However, certain critical aspects of this functionality require (often substantially) different implementations depending on properties of the notebook environment in which `davos` is used (e.g., whether the frontend is provided by Jupyter or Google Colaboratory, or which version of IPython [15] is used by the notebook kernel). To deal with this, environment-dependent parts of core features and behaviors are isolated and abstracted to “helper functions” in the `davos.implementations` subpackage. This second subpackage defines multiple, interchangeable versions of each helper function, organized into modules by the conditions that trigger their use. At runtime, `davos` detects various features in the notebook environment and selectively imports a single version of each helper function into the top-level `davos.implementations` namespace, allowing `davos.core` modules to access the proper implementations for the current notebook environment in a single, consistent location. An additional benefit of this design is that it allows maintainers, developers, and users to extend `davos` to support new, updated, or custom notebook variants by creating new `davos.implementations` modules that define their own versions of each helper function, modified from existing implementations as needed.

2.2. Software functionalities

2.2.1. The `smuggle` statement

Functionally, importing `davos` in an IPython notebook enables an additional Python keyword: “`smuggle`” (see Sec. 2.3 for details on how this works). The `smuggle` keyword-like object can be used as a drop-in replacement for Python’s built-in `import` keyword to load packages, modules, and other objects into the current namespace. However, whereas `import` will fail if the requested package is not installed locally, `smuggle` statements can handle missing packages on the fly. If a smuggled package does not exist in the local

environment, **davos** will download and install it automatically, expose its contents to Python’s **import** machinery, and load it into the namespace for immediate use.

2.2.2. The onion comment

For greater control over the behavior of **smuggle** statements, **davos** defines an additional construct called the “onion comment.” An onion comment is a special type of inline comment that may be placed on a line containing a **smuggle** statement to customize how **davos** searches for the smuggled package locally and, if necessary, downloads and installs it. Onion comments follow a simple format based on the “type comment” syntax introduced in PEP 484 [16], and are designed to make managing packages with **davos** intuitive and familiar. To construct an onion comment, users provide the name of the installer program (e.g., **pip**) and the same arguments one would use to manually install the package as desired via the command line:

```
# enable smuggle statements
import davos

# if numpy is not installed locally, pip-install it and display verbose output
smuggle numpy as np          # pip: numpy --verbose

# pip-install pandas without using or writing to the package cache
smuggle pandas as pd         # pip: pandas --no-cache-dir

# install scipy from a relative local path, in editable mode
from scipy.stats smuggle ttest_ind # pip: -e ../../pkgs/scipy
```

Occasionally, a package’s distribution name (i.e., the name used when installing it) may differ from its top-level module name (i.e., the name used when importing it). In such cases, an onion comment may be used to ensure that **davos** installs the proper package if it cannot be found locally:

```
# package is named "python-dateutil" on PyPI, but imported as "dateutil"
smuggle dateutil          # pip: python-dateutil

# package is named "scikit-learn" on PyPI, but imported as "sklearn"
from sklearn.decomposition smuggle PCA # pip: scikit-learn
```

Because onion comments may be constructed to specify any aspect of the installer’s behavior, they provide a mechanism for precisely controlling how, where, and when smuggled packages are installed. Critically, if an onion comment includes a version specifier [4], **davos** will ensure that the version of the package loaded into the notebook matches the specific version requested, or satisfies the given version constraints. If the smuggled package exists locally, **davos** will extract its version information from its metadata and compare it to the specifier provided. If the two are incompatible (or no local installation is found), **davos** will download, install, and load a suitable version of the package instead:

```
# specifically use matplotlib v3.4.2, pip-installing it if needed
smuggle matplotlib.pyplot as plt # pip: matplotlib==3.4.2

# use a version of seaborn no older than v0.9.1, but prior to v0.11
smuggle seaborn as sns          # pip: seaborn>=0.9.1,<0.11
```

118 Onion comments can also be used to smuggle specific VCS references (e.g., Git [17]
119 branches, commits, tags, etc.):

```
# use quail as the package existed on GitHub at commit 6c847a4
smuggle quail # pip: git+https://github.com/ContextLab/quail.git@6c847a4
```

120

121 **davos** processes onion comments internally before forwarding arguments to the installer
122 program. In addition to preventing onion comments from being used as a vehicle for shell
123 injection attacks, this enables **davos** to adapt its behavior based on how particular flags
124 will affect the behavior of the installer program. For example, if an onion comment con-
125 tains either the `-I/--ignore-installed`, `-U/--upgrade`, or `--force-reinstall` flag,
126 **davos** will not bother checking for a local copy of the smuggled package before installing
127 a new one:

```
# install hypertools v0.7 without first checking for it locally
smuggle hypertools as hyp # pip: hypertools==0.7 --ignore-installed

# always install the latest version of requests, including pre-releases
from requests smuggle Session # pip: requests --upgrade --pre
```

128

129 Similarly, the `--no-input` flag will temporarily enable **davos**'s non-interactive mode (see
130 Sec. 2.2.3), and installing a smuggled package into a custom directory (`<dir>`) using the
131 `--target <dir>` flag will cause **davos** to prepend `<dir>` to the module search path (i.e.,
132 `sys.path`), if necessary, so the package can be imported.

133 2.2.3. The *davos* config object

134 The **davos** config object provides a high-level interface for controlling various aspects
135 of **davos**'s behavior. After importing **davos**, the `davos.config` object (a singleton) ex-
136 poses configurable options as attributes that can be modified, displayed in the notebook,
137 or checked programmatically at runtime (see Sec. 3 for an illustrative example or Sec. 2.3
138 for implementation details and additional information). These include:

- 139 • `.active`: This attribute controls whether support for **smuggle** statements and
140 onion comments) is enabled (`True`) or disabled (`False`). When **davos** is first im-
141 ported), the `.active` attribute is set to `True`.
- 142 • `.auto_rerun`: This attribute controls how **davos** behaves when attempting to
143 **smuggle** a new version of a package that was previously imported and cannot
144 be reloaded. This can happen if the package includes extension modules that dy-
145 namically link C or C++ objects to the Python interpreter, and the code that
146 generates those objects was changed between the previously imported and to-be-
147 smuggled versions. If this attribute is set to `True`, **davos** will automatically restart
148 the notebook kernel and rerun all code up to (and including) the current **smuggle**
149 statement. If `False` (the default), **davos** will instead issue a warning, pause exe-
150 cution, and prompt the user to either restart and rerun the notebook, or continue
151 running with the previously imported package version until the next time the kernel
152 is restarted manually. Note that, as of this writing, the `.auto_rerun` attribute is
153 not supported in Google Colaboratory notebooks.

- `.confirm_install`: If `True` (default: `False`), `davos` will require user confirmation before installing a smuggled package that does not yet exist in the user's environment.
- `.noninteractive`: Setting this attribute to `True` (default: `False`) enables non-interactive mode, in which all user interactions (prompts and dialogues) are disabled. Note that in non-interactive mode, the `confirm_install` option is set to `False`. If `auto_rerun` is `False` while in non-interactive mode, `davos` will raise an exception if a smuggled package cannot be reloaded, rather than prompting the user.
- `.pip_executable`: This attribute's value specifies the path to the `pip` executable used to install smuggled packages. The default is programmatically determined from the Python environment and falls back to `sys.executable -m pip` if no executable can be found.
- `.suppress_stdout`: If this attribute is set to `True` (default: `False`), `davos` suppresses printed (console) outputs from both itself and the installer program. This can be useful when smuggling packages that need to install many dependencies and/or generate extensive output. However, if the installer program throws an error, both its stdout and stderr streams will be displayed alongside the Python traceback to allow for debugging.

The top-level `davos` namespace also defines convenience functions for setting and checking whether `davos` is active (`davos.activate()`; `davos.deactivate()`; `davos.is_active()`) as well as the `davos.configure()` function, which allows setting multiple configuration options simultaneously.

2.3. Implementation details

Although `davos` is designed to *appear* to add a new keyword to Python's vocabulary, this illusion is actually created through several "hacks" that make use of the notebook's IPython backend for processing and executing users' code. Specifically, when `davos` is first imported, or when it is activated after having been set to an inactive state, two actions are triggered: First, the `smuggle()` function is injected into the IPython user namespace. Second, the `davos` parser is registered as a custom IPython input transformer.

IPython preprocesses all executed code as plain text before it is sent to the Python compiler, in order to handle special constructs like `%magic` and `!shell` commands. `davos` uses this process to transform `smuggle` statements into syntactically valid Python code. The `davos` parser uses a regular expression to match lines of code containing `smuggle` statements (and, optionally, onion comments), extract relevant information from their text, and replace them with equivalent calls to the `smuggle()` function. For example, if a user runs a notebook cell containing

```
smuggle numpy as np      # pip: numpy>1.16,<=1.20 -vv
```

the code that is actually executed by the Python interpreter would be

```
smuggle(name="numpy", as_="np", installer="pip",
        args_str="\"numpy>1.16,<=1.20 -vv\"",
        installer_kwargs={'editable': False,
                          'spec': 'numpy>1.16,<=1.20',
                          'verbosity': 2})
```

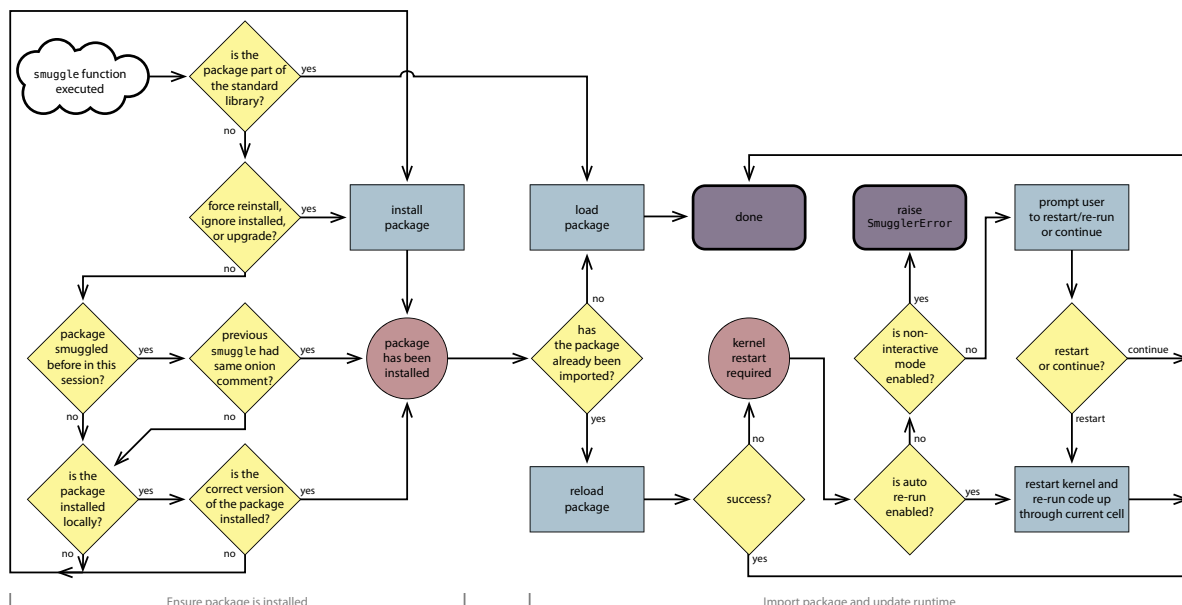



Figure 3: **smuggle() function algorithm.** At a high level, the `smuggle()` function may be conceptualized as following two basic steps. First (left), `davos` ensures that the correct version of the desired package has been installed, carrying out the installation automatically if needed. Second (right), `davos` imports the package and updates the current runtime environment.

The call to the `smuggle()` function carries out `davos`’s central logic by determining whether the smuggled package must be installed, carrying out the installation if necessary, and subsequently loading it into the namespace. This process is outlined in Figure 3. Because the `smuggle()` function is defined in the notebook namespace, it is also possible (though never necessary) to call it directly. Deactivating `davos` will delete the name “`smuggle`” from the namespace, unless its value has been overwritten and no longer refers to the `smuggle()` function. It will also deregister the `davos` parser from the set of input transformers run when each notebook cell is executed. While the overhead added by the `davos` parser is minimal, this may be useful, for example, when optimizing or precisely profiling code.

3. Illustrative Example

Across different versions of a given package, particular modules, functions, and other objects may be updated, removed, renamed, or otherwise altered in some manner. In addition to changing the behaviors of active computations, these changes can render saved objects created using one version of a package incompatible with other versions of the same package. For example, the popular `pandas` [18] library used to include the `Panel` data structure for storing 3-dimensional arrays. Since version 0.20.0, however, the `Panel` class has been deprecated, and in version 0.25.0, it was removed entirely. If a user had a dataset stored in a `Panel` object (created using an older version of `pandas`) and saved it to their disk (e.g., for later reuse or to share with other users) by serializing the `Panel` with Python’s `pickle` protocol, that dataset could not be read by any version of `pandas` from 0.25.0 or beyond. These incompatibilities are also not limited solely to traditional forms of data. For example, saved model states and other objects may reference modules, functions, attributes, classes, or other objects that may not be identical (or even present) across all versions of their associated package.

```

1  %pip install davos
2  import davos
3
4  from os.path smuggle is_file
5  smuggle joblib                                # pip: joblib<=1.2.0
6
7  davos.config.auto_rerun = True
8  smuggle numpy as np                          # pip: numpy==1.21.6
9
10 if not is_file("~/datasets/data-new.csv"):
11     smuggle pandas as pd                      # pip: pandas<0.25.0
12     tmp_data = pd.read_pickle("~/datasets/data-old.pkl")
13     tmp_data.to_frame().to_csv("~/datasets/data-new.csv")
14
15 smuggle pandas as pd                        # pip: pandas==1.3.5
16
17 davos.configure(auto_rerun=False, suppress_stdout=True, noninteractive=True)
18 smuggle tensorflow as tf                    # pip: tensorflow==2.9.2
19 from umap smuggle UMAP                     # pip: umap-learn[plot,parametric_umap]==0.5.3
20 davos.configure(suppress_stdout=False, noninteractive=False)
21
22 smuggle matplotlib.pyplot as plt           # pip: matplotlib==3.5.3
23 smuggle seaborn as sns                     # pip: seaborn==0.12.1
24 smuggle quail                               # pip: git+https://github.com/myfork/quail@6c847a4
25
26 davos.config.pip_executable = "~/envs/nb-server/bin/pip"
27 smuggle widgetsnbextension as _            # pip: widgetsnbextension==3.5.2
28 davos.config.pip_executable = "~/envs/nb-kernel/bin/pip"
29 smuggle ipywidgets                          # pip: ipywidgets==7.6.5
30
31 from tqdm.notebook smuggle tqdm             # pip: tqdm==4.62.3
32
33 data = pd.read_csv("~/datasets/data-new.csv", index_col=[0, 1])
34 smuggle sklearn                             # pip: scikit-learn<0.22.0
35 transformer = joblib.load("~/models/text-transformer.joblib")
36 smuggle sklearn                             # pip: scikit-learn==1.1.3

```

Figure 4: **Example use case for davos.** Snippets from this example are also excerpted in the main text of Section 3.

The example provided in Figure 4 demonstrates how the `davos` package can be used to circumvent these incompatibilities by carefully controlling which versions of each package are used in different parts of the notebook. The example shows how a dataset and model that require now-incompatible components of the `pandas` and `scikit-learn` [19] packages may be loaded in (using older versions of each package) and used alongside more recent versions of each package that provide new and improved functionality. When included at the top of a Jupyter notebook, the code in Figure 4 ensures that these objects will be loaded successfully and analyzed using the same set of package versions, no matter when or by whom the notebook is run.

After installing and importing `davos` (lines 1–2), we first smuggle two utilities for interacting with local files in the code below. The `smuggle` statement in line 4 loads the `is_file()` function from the Python standard library’s `os.path` module. Standard library modules are included with all Python distributions, so this line is functionally equivalent to an `import` statement and does not need or benefit from an onion comment. Line 5 loads the `joblib` package [20], installing it first, if necessary. Since `joblib`’s I/O interface has historically remained stable and backwards-compatible across releases, requiring that users have a particular exact version installed would likely be unnecessarily restrictive. However, a *future* release might introduce some breaking change. The onion comment in line 5 helps ensure the analysis notebook continues to run properly in the future by limiting allowable versions to those already released when the code was written:

```
1  %pip install davos
2  import davos
3
4  from os.path smuggle is_file
5  smuggle joblib                # pip: joblib<=1.2.0
```

Line 7 then uses the `davos.config` object to enable `davos`’s `auto_rerun` option before smuggling the next two packages: `NumPy` [21] and `pandas`. Because these packages rely heavily on custom C data types, loading the particular versions from the onion comments may require restarting the notebook kernel if different versions had been previously imported during the same interpreter session (see Sec. 2.2.3).

```
7  davos.config.auto_rerun = True
8  smuggle numpy as np          # pip: numpy==1.21.6
```

Setting the `auto_rerun` attribute to `True` is particularly useful for the installation of `pandas` in the next lines:

```
10 if not is_file("~/datasets/data-new.csv"):
11     smuggle pandas as pd          # pip: pandas<0.25.0
12     tmp_data = pd.read_pickle("~/datasets/data-old.pkl")
13     tmp_data.to_frame().to_csv("~/datasets/data-new.csv")
14
15 smuggle pandas as pd            # pip: pandas==1.3.5
```

If we suppose that the data contained in `data-old.pkl` is stored in a `Panel` object, then we must use a version of `pandas` prior to 0.25.0 (i.e., the version in which the `Panel` class was removed) to be able to load it in. Line 11 ensures that an older version of `pandas` will be imported, enabling the data to be read in (and, in line 13, saved out to a `.csv` file, which is compatible with newer versions of `pandas`).

Newer versions of **pandas** have brought substantial performance improvements, added new functionality, fixed bugs, etc. Although the original dataset had to be read in using an older version of the package, we can take advantage of those newer changes using a new **smuggle** statement and onion comment on line 15 (which specifies that version 1.3.5 should be imported). Since **pandas** had already been imported into the current workspace (on line 11), the runtime must be restarted in order to gain access to the newer version of **pandas**. The `.auto_rerun` flag set on line 7 enables this process to happen automatically, and saving out the dataset to a `.csv` file in lines 10–13 ensures that the older version of **pandas** does not need to be reinstalled.

Next, line 17 uses the `davos.configure()` function to disable the `auto_rerun` option and simultaneously enable two other options: `suppress_stdout` and `noninteractive`. With these options enabled, lines 18–19 **smuggle** **TensorFlow** [22], a powerful end-to-end platform for building and working with machine learning models, and **UMAP** [23], a package that implements a family of related manifold learning techniques. The onion comment in line 19 also specifies that **UMAP** should be installed with the optional requirements needed for its “plot” and “parametric_umap” features. Together, these two packages depend on 36 other unique packages, most of which have dependencies of their own. And if many of these are not already installed in the user’s environment, lines 18–19 could take several minutes to run. Enabling the `noninteractive` option ensures that the installation will continue automatically without user input during that time. Enabling `suppress_stdout` also suppresses console outputs from the installer that might otherwise distract from other more important outputs.

```
17  davos.configure(auto_rerun=False, suppress_stdout=True, noninteractive=True)
18  smuggle tensorflow as tf          # pip: tensorflow==2.9.2
19  from umap smuggle UMAP           # pip: umap-learn[plot,parametric_umap]==0.5.3
```

After re-enabling these two options (line 20), we next **smuggle** specific versions of three plotting packages: **Matplotlib** [24], **seaborn** [25], and **Quail** [26] (lines 22–24). Because the first two are requirements of **UMAP**’s optional “plot” feature, they will have already been installed by line 19, though possibly as different versions than those specified in the onion comments on lines 22 and 23. If the installed and specified versions are the same, these **smuggle** statements will function like standard `import` statements to load the packages into the notebook namespace. If they differ, **davos** will download the requested versions in place of the installed versions before doing so.

```
20  davos.configure(suppress_stdout=False, noninteractive=False)
21
22  smuggle matplotlib.pyplot as plt  # pip: matplotlib==3.5.3
23  smuggle seaborn as sns           # pip: seaborn==0.12.1
24  smuggle quail                    # pip: git+https://github.com/myfork/quail@6c847a4
```

Line 24 uses an onion comment to specify that **Quail** should be installed directly from a specific GitHub commit (6c847a4). This ability to install packages directly from GitHub repositories can enable developers to use forked or modified versions of other packages in their notebooks, even if those versions have not been officially released.

In lines 26–29, we demonstrate another aspect of **davos**’s functionality that supports more advanced installation scenarios. The **ipywidgets** [27] package provides an API for creating widgets, and the **widgetsnbextension** package provides the machinery needed by the notebook frontend to display them.

```

26 davos.config.pip_executable = "~/envs/nb-server/bin/pip"
27 smuggle widgetsnbextension as _ # pip: widgetsnbextension==3.5.2
28 davos.config.pip_executable = "~/envs/nb-kernel/bin/pip"
29 smuggle ipywidgets # pip: ipywidgets==7.6.5
30
31 from tqdm.notebook import tqdm # pip: tqdm==4.62.3

```

293

294 A complication is that `ipywidgets` must be installed in the same environment as the
 295 IPython kernel, whereas `widgetsnbextension` must be installed in the environment that
 296 houses the Jupyter notebook server. In standard setups, these two environments are the
 297 same. However, a common “advanced” approach entails running the notebook server from
 298 a base environment, with additional environments each providing their own separate,
 299 interchangeable IPython kernels. To accomodate this multi-environment scenario, on
 300 lines 26 and 28 we control which environments each package should be installed to. Once
 301 these two packages are installed and imported, line 30 smuggles `tqdm` [28], which display
 302 progress bars to provide status updates for running code. In Jupyter notebooks, the
 303 `tqdm.notebook` module can be imported to enable aesthetically pleasing progress bars
 304 that are displayed via `ipywidgets`, if that package is installed and importable. Therefore,
 305 to take advantage of this feature, it was important to `smuggle tqdm` after ensuring the
 306 `ipywidgets` package was available.

307 Next, we load in the reformatted dataset (line 33) and pre-trained model (line 35)
 308 that we wish to use in our analysis. In our hypothetical example, we can suppose that the
 309 model was provided as a `scikit-learn` `Pipeline` object that passes data through two
 310 pretrained models in succession. First, a trained `CountVectorizer` instance converts text
 311 data to an array of word counts. Second, the word counts are passed to a topic model [29]
 312 using a pretrained `LatentDirichletAllocation` instance.

```

33 data = pd.read_csv("~/datasets/data-new.csv", index_col=[0, 1])
34 smuggle sklearn # pip: scikit-learn<0.22.0
35 transformer = joblib.load("~/models/text-transformer.joblib")
36 smuggle sklearn # pip: scikit-learn==1.1.3

```

313

314 Let us suppose that the `Pipeline` object had been saved by its original creator using the
 315 `joblib` package, as `scikit-learn`’s documentation recommends. Because `joblib` uses
 316 the `pickle` protocol internally, the ability to save and load pre-trained models is not guar-
 317 anteed across different `scikit-learn` versions. For example, suppose that the `Pipeline`
 318 object was created using `scikit-learn` v0.21.3). In that version of `scikit-learn`, the
 319 `LatentDirichletAllocation` class is defined in `sklearn.decomposition.online_lda`.
 320 However, in version 0.22.0, that module was renamed to `_online_lda`, and in version
 321 0.22.1, it was again renamed to `_lda`.

322 In order to correctly load the model that includes the pre-trained `LatentDirichletAllocation`
 323 instance, in line 34, we first smuggle a version of `scikit-learn` prior to v0.22.0 (i.e., be-
 324 fore the first time the relevant module’s name was changed). Once the model is loaded
 325 and reconstructed in memory from a compatible package version (line 35), we upgrade to
 326 a newer version of `scikit-learn` in line 36. Taken together, the code in Figure 4 shows
 327 how `davos` can enable users to load in data and models that are incompatible with newer
 328 versions of `pandas` and `scikit-learn`, but still *analyze* and manipulate the data and
 329 model output using the latest approaches and implementations.

4. Impact

Like virtual environments, containers, and virtual machines, the **davos** package (when used in conjunction with Jupyter notebooks) provides a lightweight mechanism for sharing code and ensuring reproducibility across users and computing environments (Fig. 1). Further, **davos** enables users to fully specify (and install, as needed) any project dependencies within the same notebook. This provides a system whereby executable code (along with text and media) *and* code for setting up and configuring the project dependencies, may be combined within a single notebook file.

We designed **davos** for use in research applications. For example, in many settings, **davos** may be used as a drop-in replacement for more-difficult-to-set-up virtual environments, containers, and/or virtual machines. For researchers, this lowers barriers to sharing code. By eliminating most of the setup costs of reconstructing the original researchers’ computing environment, **davos** also lowers barriers to entry for members of the scientific community and the public who seek to run shared code.

Beyond research applications, **davos** is also useful in pedagogical settings. For example, in programming courses, instructors and students may use the **davos** package to ensure their notebooks will run correctly on others’ machines. When combined with online notebook-based platforms like Google Colaboratory, **davos** provides a convenient way to manage dependencies within a notebook, without requiring any software (beyond a web browser) to be installed on the students’ or instructors’ systems. For the same reasons, **davos** also provides an elegant means of sharing ready-to-run notebook-based demonstrations or tutorials that install their dependencies automatically.

Since its initial release, **davos** has found use in a variety of applications. In addition to managing computing environments for multiple ongoing research studies, **davos** is being used by both students and instructors in programming and methods courses such as Storytelling with Data [30] (an open course on data science, visualization, and communication) and Laboratory in Psychological Science [31] (an open course on experimental and statistical methods for psychology research) to simplify distributing lessons and submitting assignments, as well as in online demos such as **abstract2paper** [32] (an example application of GPT-Neo [33, 34]) to share ready-to-run code that installs dependencies automatically.

Our work also has several more subtle “advanced” use cases and potential impacts. Whereas Python’s built-in **import** statement is agnostic to packages’ version information, **smuggle** statements (when combined with onion comments) are version-sensitive. And because onion comments are parsed at runtime, required packages and their specified versions are installed in a just-in-time manner. Thus, it is possible in most cases to **smuggle** a specific package version or revision even if a different version has already been loaded. This enables more complex uses that take advantage of multiple versions of a package within a single interpreter session (e.g., see Sec. 3 and Fig. 4). This could be useful in cases where specific features are added or removed from a package across different versions, or in comparing the performance or functionality of particular features across different versions of the same package.

A second more subtle impact of our work is in providing a proof-of-concept of how the ability to add new “keyword-like” operators to the Python language could be specifically useful to researchers. With **davos**, we accomplish this by leveraging IPython notebooks’ internal code parsing and execution machinery. We note that, while other popular packages similarly use these mechanisms to providing notebook-specific functionality (e.g.,

[24, 35]), this approach also has the potential to be exploited for more nefarious purposes. For example, a malicious user could design a Python package that, when imported, substantially changes the notebook’s functionality by adding new *unexpected* keyword-like objects (e.g., based around common typos). We also note that this implementation approach means **davos**’s functionality is currently restricted to IPython notebook environments. However, there have been early-stage discussions of providing this sort of syntactic customizability as a core feature of the Python language, including a draft proposal [36]. In addition to enabling **davos** to be extended for use outside of notebooks, this could lead to exciting new tools that, like **davos**, extend the Python language in useful and more secure ways.

5. Conclusions

The **davos** package supports reproducible research by providing a novel lightweight system for sharing notebook-based code. But perhaps the most exciting uses of the **davos** package are those that we have *not* yet considered or imagined. We hope that the Python community will find **davos** to provide a convenient means of managing project dependencies to facilitate code sharing. We also hope that some of the more advanced applications of our package might lead to new insights or discoveries.

Author Contributions

Paxton C. Fitzpatrick: Conceptualization, Methodology, Software, Validation, Writing - Original Draft, Visualization. **Jeremy R. Manning**: Conceptualization, Resources, Validation, Writing - Review & Editing, Supervision, Funding acquisition.

Funding

Our work was supported in part by NSF grant number 2145172 to JRM. The content is solely the responsibility of the authors and does not necessarily represent the official views of our supporting organizations.

Declaration of Competing Interest

We wish to confirm that there are no known conflicts of interest associated with this publication and there has been no significant financial support for this work that could have influenced its outcome.

Acknowledgements

We acknowledge useful feedback and discussion from the students of JRM’s *Storytelling with Data* course (Winter, 2022 offering) who used preliminary versions of our package in several assignments.

References

- [1] G. van Rossum, Python reference manual, Department of Computer Science [CS] (R 9525) (1995).
- [2] Python Software Foundation, The Python Package Index (PyPI), <https://pypi.org> (2003).
- [3] conda-forge community, The conda-forge Project: Community-based Software Distribution Built on the conda Package Format and Ecosystem, <https://doi.org/10.5281/zenodo.4774217> (July 2015). doi:10.5281/zenodo.4774217.
- [4] N. Coghlan, D. Stuft, Version Identification and Dependency Specification, PEP 440, Python Software Foundation (March 2013).
- [5] B. Cannon, N. Smith, D. Stuft, Specifying Minimum Build System Requirements for Python Projects, PEP 518, Python Software Foundation (May 2016).
- [6] Anaconda, Inc., conda, <https://docs.conda.io> (2012).
- [7] S. Eustace, Poetry: Python packaging and dependency management made easy, <https://github.com/python-poetry/poetry> (December 2019).
- [8] T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier, J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov, D. Avila, S. Abdalla, C. Willing, Jupyter Notebooks – a publishing format for reproducible computational workflows, in: F. Loizides, B. Schmidt (Eds.), Positioning and Power in Academic Publishing: Players, Agents and Agendas, IOS Press, Netherlands, 2016, pp. 87–90. doi:10.3233/978-1-61499-649-1-87.
- [9] R. P. Goldberg, Survey of virtual machine research, Computer 7 (6) (1974) 34–45.
- [10] Y. Altintas, C. Brecher, M. Weck, S. Witt, Virtual Machine Tool, CIRP Annals 54 (2) (2005) 115–138. doi:[https://doi.org/10.1016/S0007-8506\(07\)60022-5](https://doi.org/10.1016/S0007-8506(07)60022-5).
- [11] M. Rosenblum, VMware’s Virtual Platform: A virtual machine monitor for commodity PCs, in: IEEE Hot Chips Symposium, IEEE, 1999, pp. 185–196.
- [12] D. Merkel, Docker: lightweight linux containers for consistent development and deployment, Linux Journal 239 (2) (2014) 2.
- [13] G. M. Kurtzer, V. Sochat, M. W. Bauer, Singularity: Scientific containers for mobility of compute, PLoS One 12 (5) (2017) e0177459.
- [14] G. R. R. Martin, A Clash of Kings, A Song of Ice and Fire, Voyager Books, 1998.
- [15] F. Pérez, B. E. Granger, IPython: a system for interactive scientific computing, Computing in science and engineering 9 (3) (2007) 21–29. doi:10.1109/MCSE.2007.53.
- [16] G. van Rossum, J. Lehtosalo, L. Langa, Type Hints, PEP 484, Python Software Foundation (September 2014).

- [17] L. Torvalds, J. Hamano, Git: Fast version control system, <https://git.kernel.org/pub/scm/git/git.git> (April 2005).
- [18] W. McKinney, Data Structures for Statistical Computing in Python, in: S. van der Walt, J. Millman (Eds.), Proceedings of the 9th Python in Science Conference, 2010, pp. 56–61. doi:10.25080/Majora-92bf1922-00a.
- [19] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, E. Duchesnay, Scikit-learn: machine learning in Python, Journal of Machine Learning Research 12 (2011) 2825–2830.
- [20] G. Varoquaux, Joblib: Computing with Python functions, <https://github.com/joblib/joblib> (July 2010).
- [21] C. R. Harris, K. J. Millman, S. J. van der Walt, R. Gommers, P. Virtanen, D. Cournapeau, E. Wieser, J. Taylor, S. Berg, N. J. Smith, R. Kern, M. Picus, S. Hoyer, M. H. van Kerkwijk, M. Brett, A. Haldane, J. F. del Río, M. Wiebe, P. Peterson, P. Gérard-Marchant, K. Sheppard, T. Reddy, W. Weckesser, H. Abbasi, C. Gohlke, T. E. Oliphant, Array programming with NumPy, Nature 585 (7825) (2020) 357–362. doi:10.1038/s41586-020-2649-2.
- [22] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Jozefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu, X. Zheng, TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems, software available from tensorflow.org (2015). URL <https://www.tensorflow.org/>
- [23] L. McInnes, J. Healy, N. Saul, L. Großberger, UMAP: Uniform Manifold Approximation and Projection, Journal of Open Source Software 3 (29) (2018) 861. doi:<https://doi.org/10.21105/joss.00861>.
- [24] J. D. Hunter, Matplotlib: A 2D graphics environment, Computing in Science and Engineering 9 (3) (2007) 90–95. doi:10.1109/MCSE.2007.55.
- [25] M. L. Waskom, seaborn: statistical data visualization, Journal of Open Source Software 6 (60) (2021) 3021. doi:10.21105/joss.03021.
- [26] A. C. Heusser, P. C. Fitzpatrick, C. E. Field, K. Ziman, J. R. Manning, Quail: a Python toolbox for analyzing and plotting free recall data, Journal of Open Source Software 10.21105/joss.00424 (2017).
- [27] J. Frederic, J. Grout, Jupyter Widgets Contributors, ipywidgets: Interactive Widgets for the Jupyter Notebook, <https://github.com/jupyter-widgets/ipywidgets> (August 2015).
- [28] C. da Costa-Luis, S. K. Larroque, K. Altendorf, H. Mary, richardsheridan, M. Korobov, N. Raphael, I. Ivanov, M. Bargull, N. Rodrigues, G. Chen, A. Lee, C. Newey,

CrazyPython, JC, M. Zugnoni, M. D. Pagel, mjstevens777, M. Dektyarev, A. Rothberg, A. Plavin, D. Panteleit, F. Dill, FichteFoll, G. Sturm, HeoHeo, H. van Kemenade, J. McCracken, MapleCCC, M. Nordlund, tqdm: A Fast, Extensible Progress Bar for Python and CLI, <https://github.com/tqdm/tqdm> (September 2022). doi:10.5281/zenodo.595120.

[29] D. M. Blei, A. Y. Ng, M. I. Jordan, Latent dirichlet allocation, Journal of Machine Learning Research 3 (2003) 993–1022.

[30] J. R. Manning, Storytelling with Data, <https://github.com/ContextLab/storytelling-with-data> (June 2021). doi:10.5281/zenodo.5182775.

[31] J. Manning, ContextLab/experimental-psychology: v1.0 (Spring, 2022), <https://github.com/ContextLab/experimental-psychology/tree/v1.0> (May 2022). doi:10.5281/zenodo.6596762.

[32] J. R. Manning, abstract2paper, <https://github.com/ContextLab/abstract2paper> (June 2021). doi:10.5281/zenodo.7261831.

[33] L. Gao, S. Biderman, S. Black, L. Golding, T. Hoppe, C. Foster, J. Phang, H. He, A. Thite, N. Nabeshima, S. Presser, C. Leahy, The Pile: An 800GB Dataset of Diverse Text for Language Modeling, arXiv preprint arXiv:2101.00027 (2020).

[34] S. Black, L. Gao, P. Wang, C. Leahy, S. Biderman, GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow, <http://github.com/eleutherai/gpt-neo> (2021).

[35] A. C. Heusser, K. Ziman, L. L. W. Owen, J. R. Manning, HyperTools: a Python toolbox for gaining geometric insights into high-dimensional data, Journal of Machine Learning Research 18 (152) (2018) 1–6.

[36] M. Shannon, Syntactic Macros, Draft PEP 638, Python Software Foundation (September 2020).