

davos: a Python package “smuggler” for constructing lightweight reproducible notebooks

Paxton C. Fitzpatrick, Jeremy R. Manning*

*Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755*

Abstract

Reproducibility is a core requirement of modern scientific research. For computational research, reproducibility means that code should produce the same results, even when run on different systems. A standard approach to ensuring reproducibility entails packaging a project’s dependencies along with its primary code base. Existing solutions vary in how deeply these dependencies are specified, ranging from virtual environments (which specify all Python package versions), to containers (which also specify the operating system), to virtual machines (which also specify hardware layers of the system). Each of these existing solutions requires installing or setting up a system for running the desired code that must be packaged alongside the primary code base. Here we propose a lighter-weight solution than virtual environments: the **davos** library. When used in combination with a notebook-based Python project, the **davos** library provides a mechanism for specifying (and automatically installing) the correct package versions of the project’s dependencies. This enables researchers to share a complete reproducible environment using a single Jupyter notebook file.

Keywords: Reproducibility, Open science, Python, Jupyter Notebook, Google Colaboratory, Package management

*Corresponding author

Email address: `Jeremy.R.Manning@Dartmouth.edu` (Jeremy R. Manning)

Required Metadata

Current code version

Nr.	Code metadata description	Metadata value
C1	Current code version	v0.1.1
C2	Permanent link to code/repository used for this code version	https://github.com/ContextLab/davos/tree/v0.1.1
C3	Code Ocean compute capsule	
C4	Legal Code License	MIT
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	Python, JavaScript, PyPI/pip, IPython, Jupyter, Ipykernel, PyZMQ. Additional tools used for tests: pytest, Selenium, Requests, mypy, GitHub Actions
C7	Compilation requirements, operating environments, and dependencies	Dependencies: Python ≥ 3.6 , packaging, setuptools. Supported OSes: MacOS, Linux, Unix-like. Supported IPython environments: Jupyter notebooks, JupyterLab, Google Colaboratory, Binder, IDE-based notebook editors.
C8	Link to developer documentation/manual	https://github.com/ContextLab/davos#readme
C9	Support email for questions	contextualdynamics@gmail.com

Table 1: Code metadata

1. Motivation and significance

Code sharing is a core component of the open science movement that has inspired the development of a full ecosystem of tools and packages. However, sharing code, in and of itself, does not guarantee that others will be able to *reproduce* the desired results. For example, research code often requires installing other software packages that extend the implementation language’s basic functionality. Within the Python community [1], external packages that are published in the most popular repositories [2, 3] are associated with version numbers and tags that enable users to guarantee that they are installing exactly the same code across different computing environments. Despite that it is *possible* to manually install the intended version numbers of every dependency of a Python script or package, doing so may cause conflicts within the

13 user’s computing environment that interfere with the functionality of *other*
14 code.

15 To facilitate code sharing, the Python community has developed a broad
16 set of approaches and tools (Fig. 1). At one extreme, simply publishing a set
17 of Python scripts (.py files) may enable others to use or gain insights into
18 the relevant work. Because Python is installed by default on most modern
19 operating systems, for some projects this may be sufficient. Another popu-
20 lar approach entails creating JSON files, called Jupyter notebooks [4], that
21 comprise a mix of text, executable code, and embedded media. Notebooks
22 may call or import external scripts or libraries in order to provide a more
23 compact and readable experience for users. Each of these systems (Python
24 scripts and notebooks) provides a convenient means of sharing code, with the
25 caveat that they do not specify the computing environment in which the code
26 is executed. Therefore the functionality of code shared using these systems
27 cannot be guaranteed across different computing environments.

28 At another extreme, virtual machines [5, 6, 7] provide a hardware-level
29 simulation of the desired system. Virtual machines are typically isolated
30 from the user’s system, such that installing or running software on a virtual
31 machine does not impact the user’s primary operating system or computing
32 environment. Containers [e.g., 8, 9] provide a similar “isolated” experience.
33 Although containerized environments do not specify hardware-level opera-
34 tions, they are typically packaged with a complete operating system, in ad-
35 dition to a complete copy of Python and any relevant package dependencies.
36 Virtual environments [e.g., 10] also provide a computing environment that
37 is largely separated from the user’s main environment. They incorporate
38 a copy of Python and the target software’s dependencies, but virtual envi-
39 ronments do not specify or reproduce an operating system for the runtime
40 environment. Each of these systems (virtual machines, containers, and vir-
41 tual environments) guarantees (to differing degrees– at the hardware level,
42 operating system level, and Python environment level, respectively) that the
43 relevant code will run similarly for different users. However, each of these
44 systems also relies on additional software that can be resource intensive or
45 burdensome to install or configure.

46 We designed **davos** to occupy a “sweet spot” between these extremes.
47 **davos** is a notebook-installable package that adds functionality to the default
48 notebook experience. Like standard Jupyter notebooks, **davos**-enhanced
49 notebooks allows researchers to include text, executable code, and media
50 within a single file. No further setup or installation is required, beyond what
51 is needed to run standard Jupyter notebooks. And like virtual environments
52 **davos** provides a convenient mechanism for fully specifying (and installing, as
53 needed) a complete set of Python dependencies, including package versions.

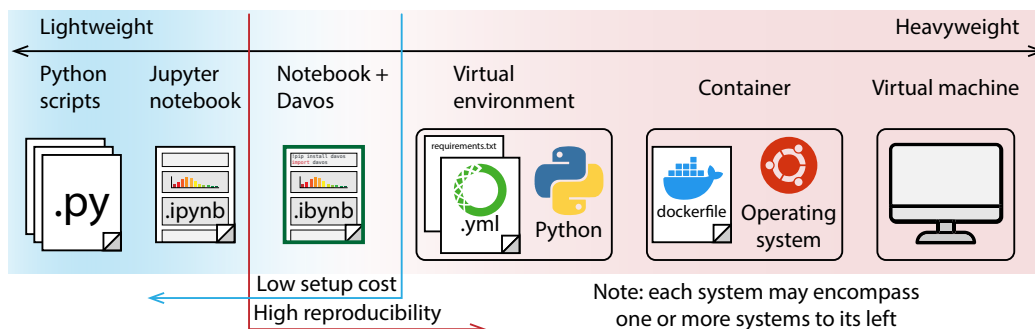


Figure 1: **Systems for sharing code within the Python ecosystem.** From left to right: plain-text **Python scripts** (`.py` files) provide the most basic “system” for sharing raw code. Scripts may reference external libraries, but those libraries are not automatically installed on other users’ systems, nor is any version checking performed by default. **Jupyter notebooks** (`.ipynb` files) comprise embedded text, executable code, and media (including rendered figures, code output, etc.). When the **davos** library is imported into a Jupyter notebook, the notebook’s functionality is extended to automatically install the required external libraries (at their correct versions, when specified). **Virtual environments** install an isolated copy of Python and all required dependencies. This typically requires defining a `requirements.txt` file that lists all dependencies (including version numbers) along with an environment (`.yml`) file that specifies how the virtual environment should be configured. **Containers** provide a means of defining an isolated environment that includes a complete operating system (independent of the user’s operating system), in addition to a virtual environment or other configurations needed to provide the necessary computing environment. Containers are typically defined using specification files (e.g., a plain-text `dockerfile`) that instruct the virtualization engine regarding how to build the virtual environment. **Virtual machines** provide a complete hardware-level simulation of the computing environment. In addition to simulating specific hardware, virtual machines (typically specified using binary images files) must also define operating system-level properties of the computing environment. Systems to the left of the blue vertical line entail sharing individual files, with no additional installation or configuration needed to run the target software. Systems to the right of the red vertical line provide high reproducibility by supporting precise control over dependencies and versioning. Notebooks enhanced using the **davos** library are easily shareable and require minimal setup costs, while also facilitating high reproducibility by enabling precise control over project dependencies.

54 2. Software description

55 2.1. Software architecture

56 The `davos` package consists of two interdependent subpackages. The
57 first, `davos.core`, comprises a set of modules that implement the bulk of the
58 package’s core functionality, including pipelines for installing and validating
59 packages, custom parsers for the `smuggle` statement (see Section 2.2.1) and
60 onion comment (see Section 2.2.2), and a runtime interface for configuring
61 `davos`’s behavior (see Section 2.2.3). However, certain critical aspects of
62 this functionality require (often substantially) different implementation ap-
63 proaches depending on various properties of the notebook environment in
64 which `davos` is used (e.g., whether the frontend is provided by Jupyter or
65 Google Colaboratory, or which version of IPython [11] is used by the note-
66 book kernel). To deal with this, environment-dependent parts of core fea-
67 tures and behaviors are isolated and abstracted to “helper functions” in the
68 `davos.implementations` subpackage. This second subpackage defines multi-
69 ple, interchangeable versions of each helper function, organized into modules
70 by the conditions that trigger their use. At runtime, `davos` detects various
71 features in the notebook environment and selectively imports a single version
72 of each helper function into the top-level `davos.implementations` names-
73 pace, allowing `davos.core` modules to access the correct implementations
74 for the current notebook environment from a single, constant location. An
75 additional benefit of this design pattern is that it makes adding support for
76 new or updated notebook variants to `davos` in the future relatively easy.

77 2.2. Software functionalities

78 2.2.1. The `smuggle` statement

79 Importing `davos` in a Jupyter notebook enables an additional Python
80 keyword: “`smuggle`”. The `smuggle` statement can be used as a drop-in re-
81 placement for Python’s built-in `import` statement to load libraries, modules,
82 and other objects into the current namespace. However, whereas `import` will
83 fail if the requested package is not installed locally, `smuggle` statements can
84 handle missing packages on the fly. If a smuggled package does not exist in
85 the local environment, `davos` will install it, expose its contents to Python’s
86 `import` machinery, and load it into the namespace for immediate use.

87 2.2.2. The onion comment

88 For greater control over the behavior of `smuggle` statements, `davos` de-
89 fines an additional construct called the “onion comment”. An onion comment
90 is a special type of inline comment that may be placed on a line containing a

91 `smuggle` statement to customize how `davos` searches for the smuggled pack-
92 age locally and, if necessary, downloads and installs it. Onion comments
93 follow a simple syntax based on the “type comment” syntax introduced in
94 PEP 484 [10], and are designed to make managing packages with `davos` intu-
95 itive and familiar. To construct an onion comment, simply provide the name
96 of the installer program (e.g., `pip`) and the same arguments one would use
97 to manually install the package as desired via the command line (see Fig. 2,
98 lines 3–10 for examples).

99 Onion comments are useful when smuggling a package whose distribution
100 name (i.e., the name used when installing it) is different from its top-level
101 module name (i.e., the name used when importing it; e.g., Fig. 2, lines 12–13).
102 However, the most powerful use of the onion comment is making `smuggle`
103 statements *version-sensitive*. If an onion comment includes a version specifier
104 [12] (e.g., Fig. 2, lines 15–19), `davos` will ensure that the version of the
105 package loaded into the notebook matches the specific version requested, or
106 satisfies the given version constraints. If the smuggled package exists locally,
107 `davos` will extract its version info from its metadata and compare it to the
108 specifier provided. If the two are incompatible (or no local installation is
109 found), `davos` will install and load a suitable version of the package instead.
110 Onion comments can similarly be used to smuggle specific VCS references
111 (e.g., Git [13] branches, commits, tags, etc.; Fig. 2, lines 21–22).

112 `davos` processes onion comments internally before forwarding arguments
113 to the installer program. In addition to preventing onion comments from
114 being used as a vehicle for shell injection attacks, this allows `davos` take cer-
115 tain logical actions when particular arguments are passed (e.g., Fig. 2, lines
116 24–28). For example, the `--force-reinstall`, `-I/--ignore-installed`,
117 and `-U/--upgrade` flag will all cause `davos` to skip searching for a smug-
118 gled package locally before installing a new copy; `--no-input` will disable
119 `davos`’s input prompts in addition to the installer program’s; and installing
120 a package into `<dir>` with `--target <dir>` will cause `dir` to be prepended
121 to the module search path (`sys.path`), if necessary, so the package can be
122 imported.

123 *2.2.3. The `davos` config*

124 *2.2.4. Additional functionality*

125 3. Illustrative Examples

126 4. Impact

127 Like virtual environments, containers, and virtual machines, the `davos` li-
128 brary (when used in conjunction with Jupyter notebooks) provides a lightweight
129 mechanism for sharing code and ensuring reproducibility across users and

```

1  import davos
2
3  # if numpy is not installed locally, pip-install it and display verbose output
4  smuggle numpy as np    # pip: numpy --verbose
5
6  # pip-install pandas without using or writing to the package cache
7  smuggle pandas as pd   # pip: pandas --no-cache-dir
8
9  # install scipy from a relative local path, in editable mode
10 from scipy.stats smuggle ttest_ind    # pip: -e ../../pkgs/scipy
11
12 smuggle dateutil        # pip: python-dateutil
13 from sklearn.decomposition smuggle PCA    # pip: scikit-learn
14
15 # specifically use matplotlib v3.4.2, pip-installing it if needed
16 smuggle matplotlib.pyplot as plt    # pip: matplotlib==3.4.2
17
18 # use a version of seaborn no older than v0.9.1, but before v0.11
19 smuggle seaborn as sns    # pip: seaborn>=0.9.1,<0.11
20
21 # use quail as the package existed on GitHub at commit 6c847a4
22 smuggle quail    # pip: git+https://github.com/ContextLab/quail.git@6c847a4
23
24 # install hypertools v0.7 without first checking for it locally
25 smuggle hypertools as hyp    # pip: hypertools==0.7 --ignore-installed
26
27 # always install the latest version of requests, including pre-releases
28 from requests smuggle Session    # pip: requests --upgrade --pre

```

Figure 2: Example `smuggle` statements and accompanying onion comments.

130 computing environments (Fig. 1). Further, `davos` enables users to fully
 131 specify (and install, as needed) any project dependencies within the same
 132 notebook. This provides a system whereby executable code (along with text
 133 and media) *and* code for setting up and configuring the project dependencies,
 134 may be combined within a single notebook file.

135 We designed `davos` for use in research applications. For example, in many
 136 settings `davos` may be used as a drop-in replacement for more-difficult-to-
 137 set-up virtual environments, containers, and/or virtual machines. For re-
 138 searchers, this lowers barriers to sharing code. By eliminating most of the
 139 setup costs of reconstructing the original researchers’ computing environ-
 140 ment, `davos` also lowers barriers to entry for members of the scientific com-
 141 munity and the public who seek to *benefit* from shared code.

142 Beyond research applications, **davos** is also useful in pedagogical settings.
143 For example, in programming courses, instructors and students may import
144 the **davos** library into their notebooks to provide a simple means of ensur-
145 ing their code will run on others’ machines. When combined with online
146 notebook-based platforms like Google Colaboratory, **davos** provides a con-
147 venient way to manage dependencies within a notebook, without requiring
148 any software (beyond a web browser) to be installed on the students’ or in-
149 structors’ systems. For the same reasons, **davos** also provides an elegant
150 means of sharing ready-to-run notebook-based demonstrations that install
151 their dependencies automatically.

152 Our work also has several more subtle “advanced” use cases and poten-
153 tial impacts. Whereas Python’s built-in **import** statement is agnostic to
154 packages’ version numbers, **smuggle** statements (when combined with onion
155 comments) are version-sensitive. This enables multiple versions of a single li-
156 brary to be imported within the same notebook. This could be useful in cases
157 where specific features were added or removed from a package across differ-
158 ent versions, or in comparing the performance or functionality of particular
159 features across different versions of the same package.

160 A second advanced use case is in providing a proof-of-concept of how one
161 can add new “keywords” to the Python language by leveraging the error-
162 handling mechanisms. This could lead to exciting new tools that, like **davos**,
163 extend the Python language in useful ways. We note that our approach
164 to adding the **smuggle** keyword to Python when **davos** is imported into a
165 notebook-based environment also has the potential to be exploited for more
166 nefarious purposes. For example, a malicious user could use a similar ap-
167 proach (e.g., in a different library) to substantially change a notebook’s func-
168 tionality by adding new *unexpected* keyword-like objects (e.g., based around
169 common typos). This could lead to difficult-to-predict changes in a note-
170 book’s behavior once the malicious library was imported. This highlights an
171 important reason why security-conscious users would be well-served to only
172 make use of libraries from trusted sources, or whose code is publicly available
173 for review.

174 5. Conclusions

175 The **davos** library supports reproducible research by providing a novel
176 lightweight system for sharing notebook-based code. But perhaps the most
177 exciting uses of the **davos** library are those that we have *not* yet considered
178 or imagined. We hope that the Python community will find **davos** to pro-
179 vide a convenient means of managing project dependencies to facilitate code

180 sharing. We also hope that some of the more advanced applications of our
181 library might lead to new insights or discoveries.

182 **Author Contributions**

183 Conceptualization: PCF and JRM. Methodology: PCF and JRM. Imple-
184 mentation: PCF. Validation: PCF. Testing: PCF and JRM. Writing: PCF
185 and JRM.

186 **Funding**

187 Our work was supported in part by NSF grant number 2145172 to JRM.
188 The content is solely the responsibility of the authors and does not necessarily
189 represent the official views of our supporting organizations.

190 **Declaration of Competing Interest**

191 We wish to confirm that there are no known conflicts of interest associated
192 with this publication and there has been no significant financial support for
193 this work that could have influenced its outcome.

194 **Acknowledgements**

195 We acknowledge useful feedback and discussion from the students of
196 JRM’s *Storytelling with Data* course (Winter, 2022 offering) who used pre-
197 liminary versions of our library in several assignments.

198 **References**

- 199 [1] G. van Rossum, Python reference manual, Department of Computer
200 Science [CS] (R 9525) (1995).
- 201 [2] Python Software Foundation, The Python Package Index (PyPI),
202 <https://pypi.org> (2003).
- 203 [3] Anaconda, Inc., conda, <https://docs.conda.io> (2012).
- 204 [4] T. Kluyver, B. Ragan-Kelley, F. Pérez, B. Granger, M. Bussonnier,
205 J. Frederic, K. Kelley, J. Hamrick, J. Grout, S. Corlay, P. Ivanov,
206 D. Avila, S. Abdalla, C. Willing, Jupyter notebooks – a publish-
207 ing format for reproducible computational workflows., in: F. Loizides,
208 B. Schmidt (Eds.), *Positioning and Power in Academic Publishing: Play-*
209 *ers, Agents and Agendas*, IOS Press, Netherlands, 2016, pp. 97–90.
210 doi:10.3233/978-1-61499-649-1-87.

- 211 [5] R. P. Goldberg, Survey of virtual machine research, *Computer* 7 (6)
212 (1974) 34–45.
- 213 [6] Y. Altintas, C. Brecher, M. Weck, S. Witt, [Virtual ma-](#)
214 [chine tool](#), *CIRP Annals* 54 (2) (2005) 115–138. doi:[https:](https://doi.org/10.1016/S0007-8506(07)60022-5)
215 [//doi.org/10.1016/S0007-8506\(07\)60022-5](https://doi.org/10.1016/S0007-8506(07)60022-5).
216 URL [https://www.sciencedirect.com/science/article/pii/](https://www.sciencedirect.com/science/article/pii/S0007850607600225)
217 [S0007850607600225](https://www.sciencedirect.com/science/article/pii/S0007850607600225)
- 218 [7] M. Rosenblum, VMware’s Virtual Platform: A virtual machine monitor
219 for commodity PCs, in: *IEEE Hot Chips Symposium*, IEEE, 1999, pp.
220 185–196.
- 221 [8] D. Merkel, Docker: lightweight linux containers for consistent develop-
222 ment and deployment, *Linux Journal* 239 (2) (2014) 2.
- 223 [9] G. M. Kurtzer, V. Sochat, M. W. Bauer, Singularity: Scientific contain-
224 ers for mobility of compute, *PLoS One* 12 (5) (2017) e0177459.
- 225 [10] G. van Rossum, J. Lehtosalo, L. Langa, [Type Hints](#), PEP 484, Python
226 Software Foundation (September 2014).
227 URL <https://www.python.org/dev/peps/pep-0484>
- 228 [11] F. Pérez, B. E. Granger, IPython: a system for interactive scientific
229 computing, *Computing in science & engineering* 9 (3) (2007) 21–29.
230 doi:[10.1109/MCSE.2007.53](https://doi.org/10.1109/MCSE.2007.53).
- 231 [12] N. Coghlan, D. Stufft, [Version identification and dependency specifica-](#)
232 [tion](#), PEP 440, Python Software Foundation (March 2013).
233 URL <https://peps.python.org/pep-0440>
- 234 [13] L. Torvalds, J. Hamano, Git: Fast version control system, [https://](https://git.kernel.org/pub/scm/git/git.git)
235 git.kernel.org/pub/scm/git/git.git (April 2005).